

RISC 5-stage Pipeline and Extension

Lecture 3

September 13th, 2023

Jae W. Lee (jaewlee@snu.ac.kr)

Computer Science and Engineering

Seoul National University

Slide credits: Some slides taken from UC Berkeley CS 152, MIT 6.823, and instructor's slides from Elsevier Inc.

Review – CPU Performance

$$\text{CPU time} = \frac{\text{Seconds}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Cycles}}{\text{Instruction}} \times \frac{\text{Seconds}}{\text{Cycle}}$$

■ Performance depends on

- Algorithm
- Programming language
- Compiler
- Instruction set architecture (ISA)
- Core organization (also called microarchitecture)
- Fabrication technology

} **Software stack**

- **Architecture**

} **Implementation**

Review – Determinates of CPU Performance

■ $\text{CPU time} = \text{Instruction_count} \times \text{CPI} \times \text{clock_cycle}$

	Instruction_ count	CPI	clock_cycle
Algorithm			
Programming language			
Compiler			
ISA			
Core organization			
Technology			

Review – Determinates of CPU Performance

■ $\text{CPU time} = \text{Instruction_count} \times \text{CPI} \times \text{clock_cycle}$

	Instruction_ count	CPI	clock_cycle
Algorithm	X	X	
Programming language	X	X	
Compiler	X	X	
ISA	X	X	X
Core organization		X	X
Technology			X

Review – RISC-V ISA

■ 32-bit fixed format instruction (6 formats)

Name (Field Size)	Field						Comments
	7 bits	5 bits	5 bits	3 bits	5 bits	7 bits	
R-type	funct7	rs2	rs1	funct3	rd	opcode	Arithmetic instruction format
I-type	immediate[11:0]		rs1	funct3	rd	opcode	Loads & immediate arithmetic
S-type	immed[11:5]	rs2	rs1	funct3	immed[4:0]	opcode	Stores
SB-type	immed[12,10:5]	rs2	rs1	funct3	immed[4:1,11]	opcode	Conditional branch format
UJ-type	immediate[20,10:1,11,19:12]				rd	opcode	Unconditional jump format
U-type	immediate[31:12]				rd	opcode	Upper immediate format

■ 32 general-purpose registers (GPRs)

- x0 is hardwired to zero (read-only)

■ 3-operand, reg-reg arithmetic instruction

- e.g., `add x28, x5, x6` # $x28 = x5 + x6$

■ Single address mode for load/store: base + displacement

- no indirection


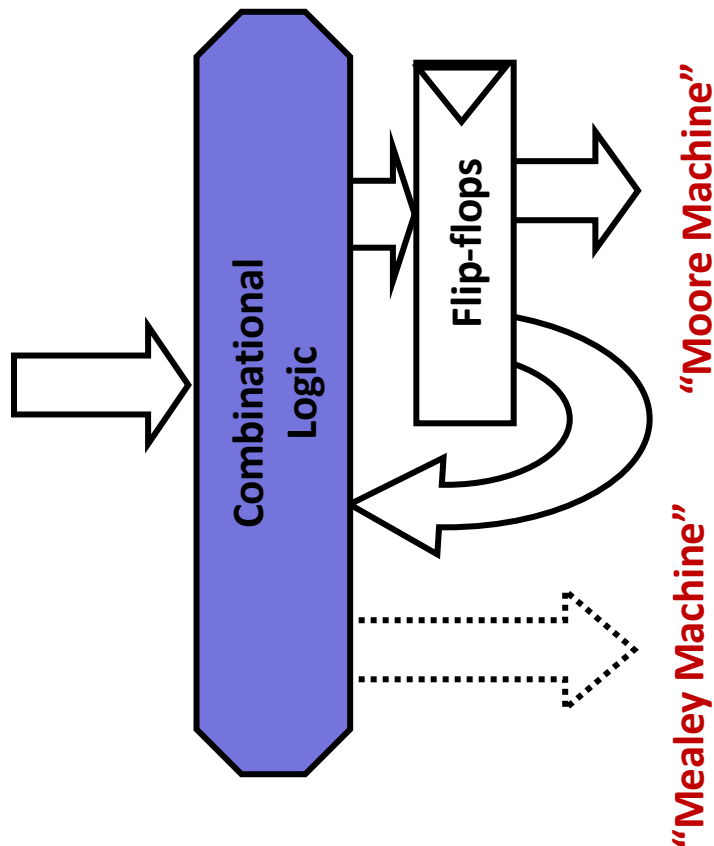
Outline

Textbook: Appendix C.1-C.5

- **ISA Implementation Basics**
- **Classic 5-stage Pipeline for RISC: A First Shot**
- **Pipeline Hazards**
 - Structural Hazards
 - Data Hazards
 - Control Hazards
- **Exception/Interrupt Handling**
- **Handling Multi-Cycle Operations**

ISA Implementation Basics *Synchronous + Digital*

- Finite State Machines = Combinational logic + Flip-Flops


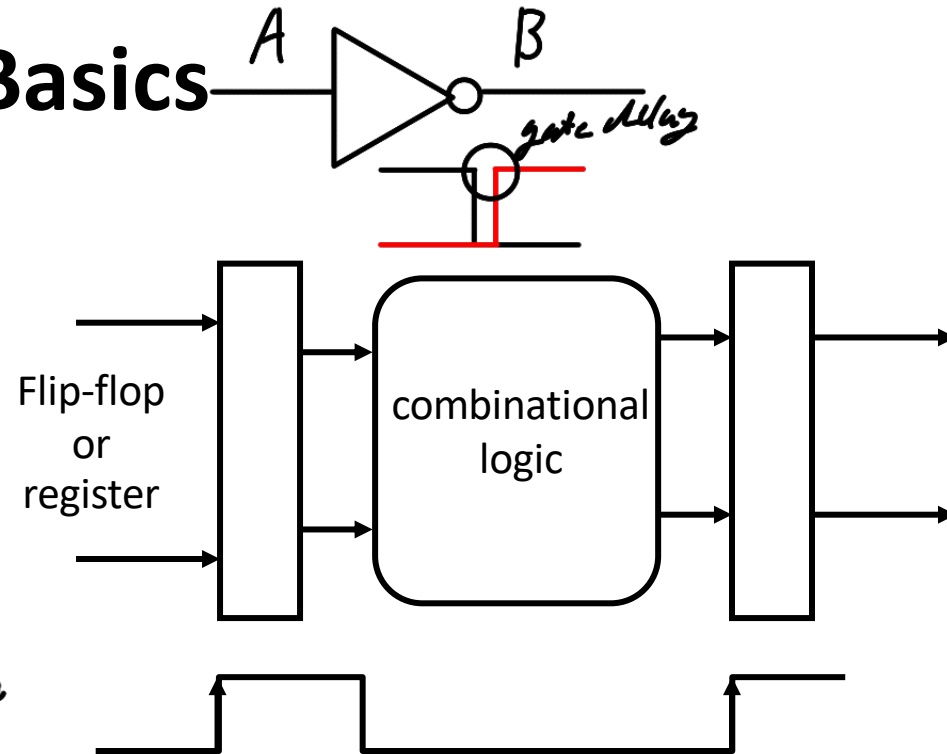


Input	State _{old}	State _{new}	Div
0	00	00	0
0	01	10	0
0	10	01	1
1	00	01	0
1	01	00	1
1	10	10	1

ISA Implementation Basics

■ What's a Clock Cycle?

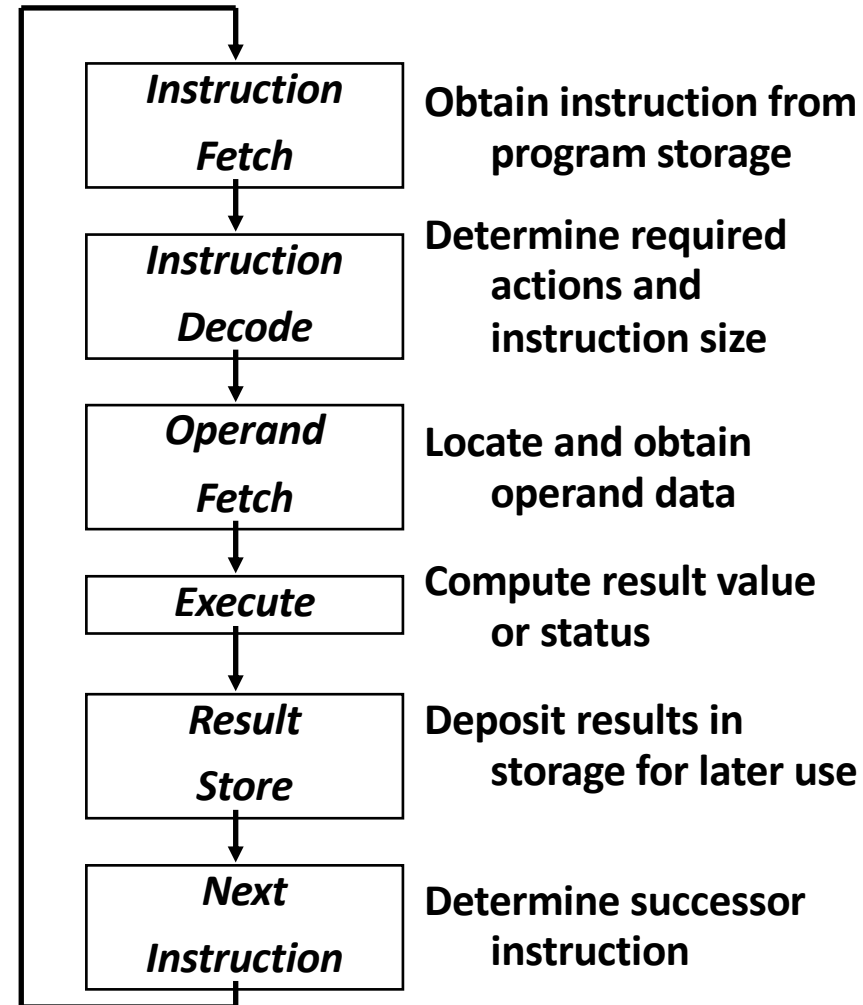

- Old days: 10 levels of gates
- Today: determined by numerous time-of-flight issues + gate delays + *Wire*
 - clock propagation, wire



Source: A. Danowitz et al., CPU DB: Recording microprocessor history, ACM Queue, 2012.

ISA Implementation Basics

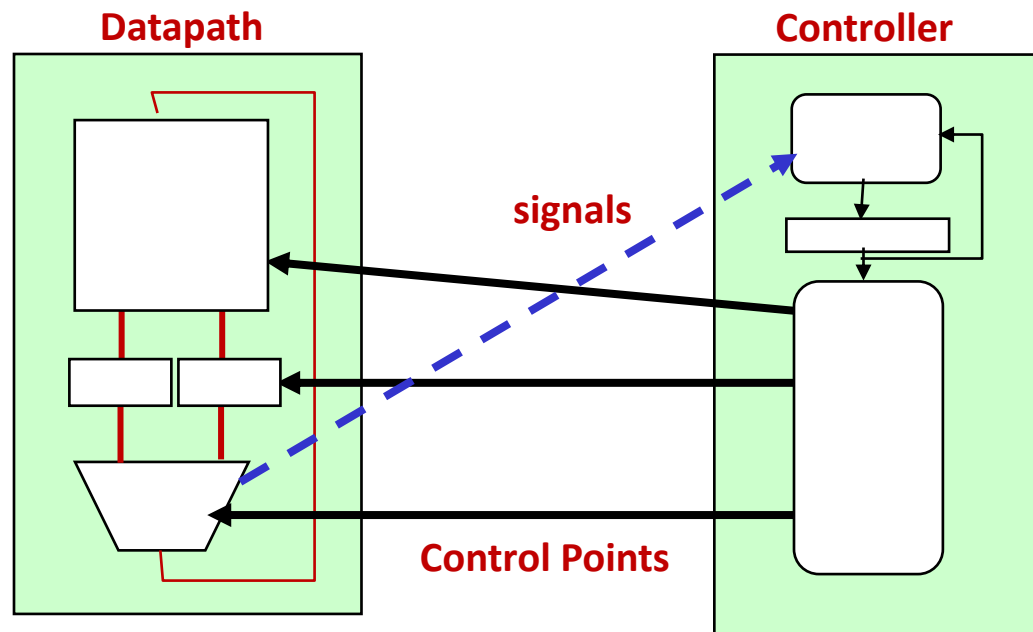
■ Fundamental Execution Cycle



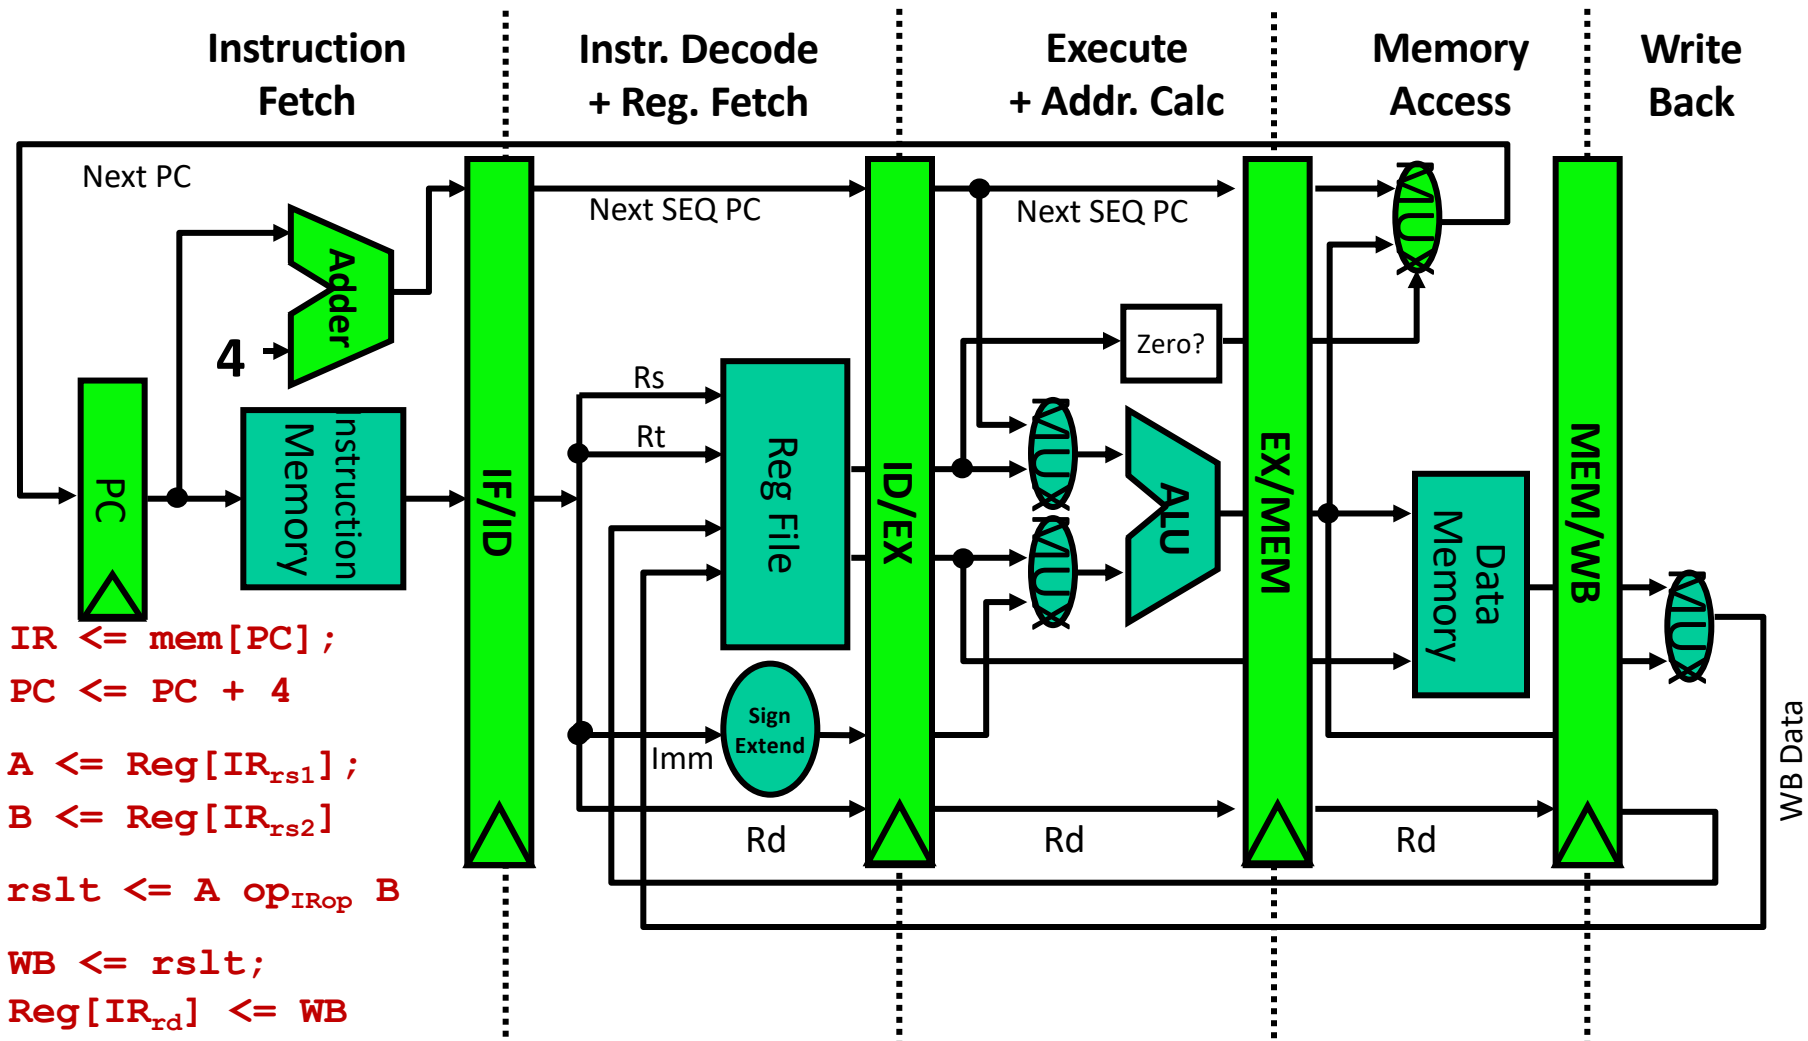
ISA Implementation Basics

■ Datapath vs Control

- Datapath: Storage, FU, interconnect sufficient to perform desired functions
 - Inputs are Control Points
 - Outputs are signals
- Controller: State machine to orchestrate operation on data path
 - Based on desired function and signals

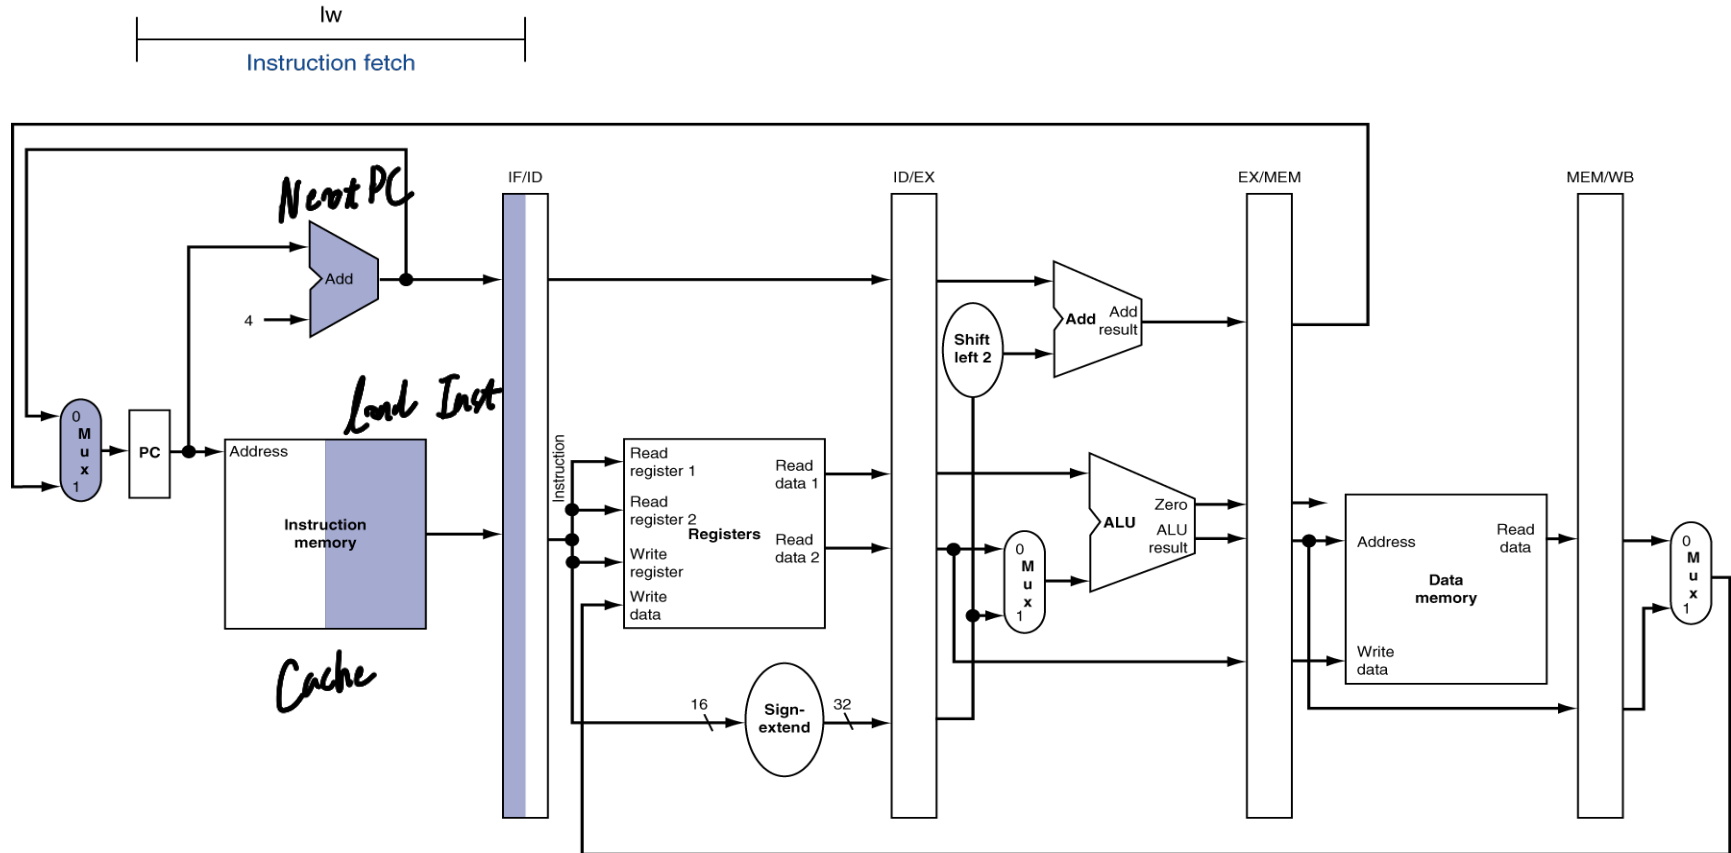


Classic 5-stage Pipeline for RISC: A First Shot



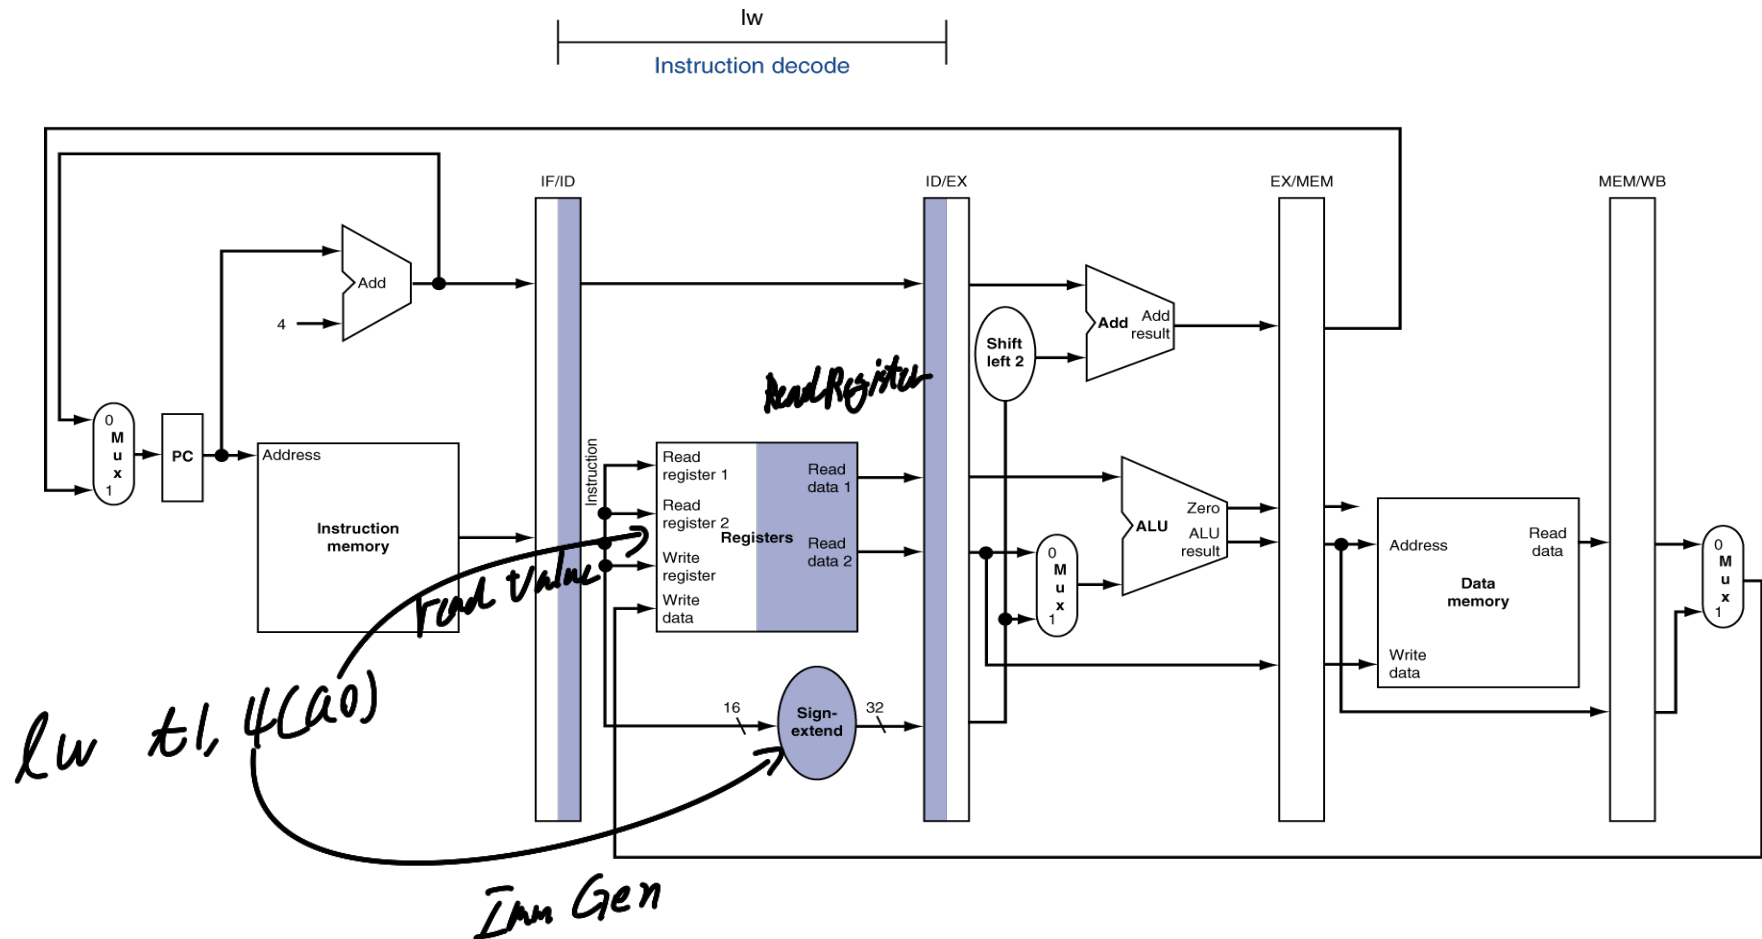
Classic 5-stage Pipeline for RISC: A First Shot

■ Example: executing load (lw) instruction – IF stage



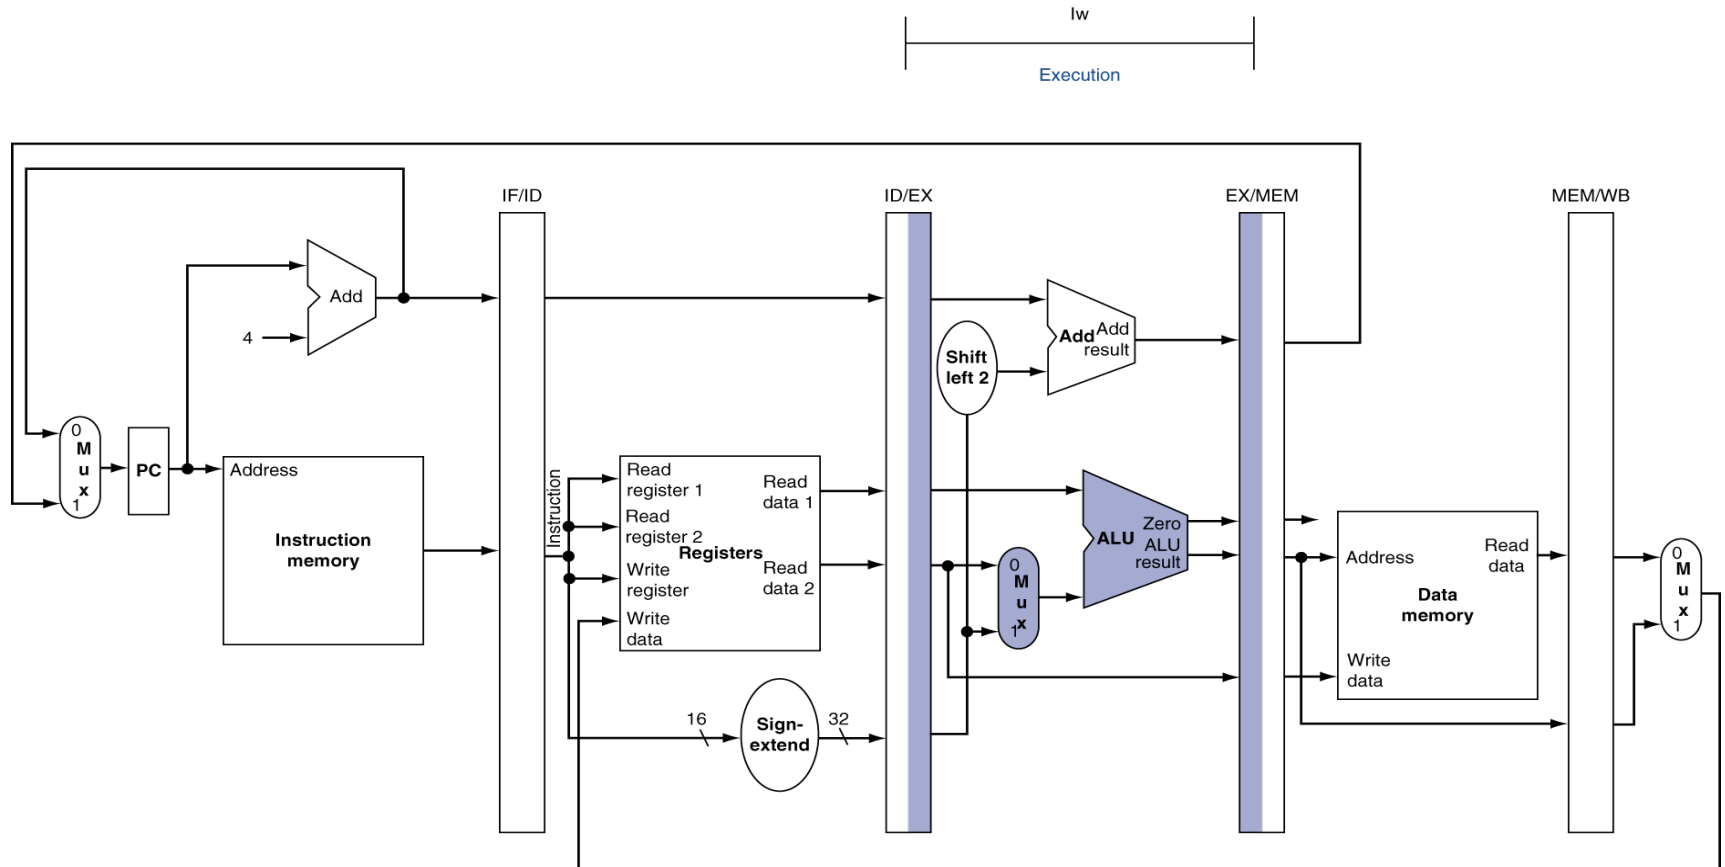
Classic 5-stage Pipeline for RISC: A First Shot

- Example: executing load (lw) instruction – **ID stage**



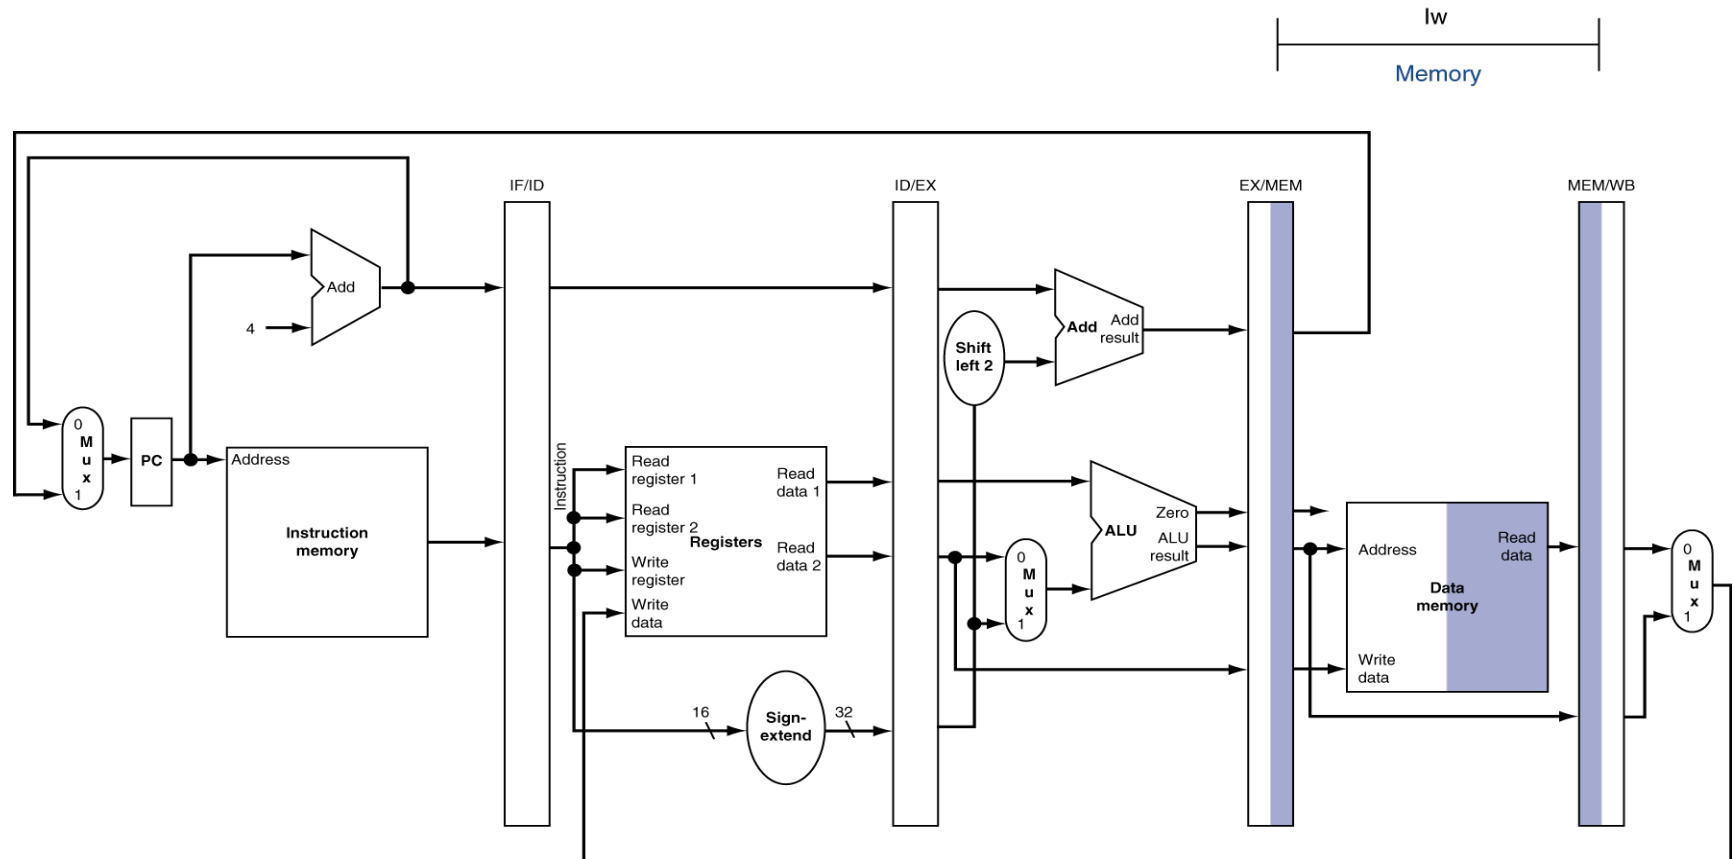
Classic 5-stage Pipeline for RISC: A First Shot

- Example: executing load (lw) instruction – **EX stage**



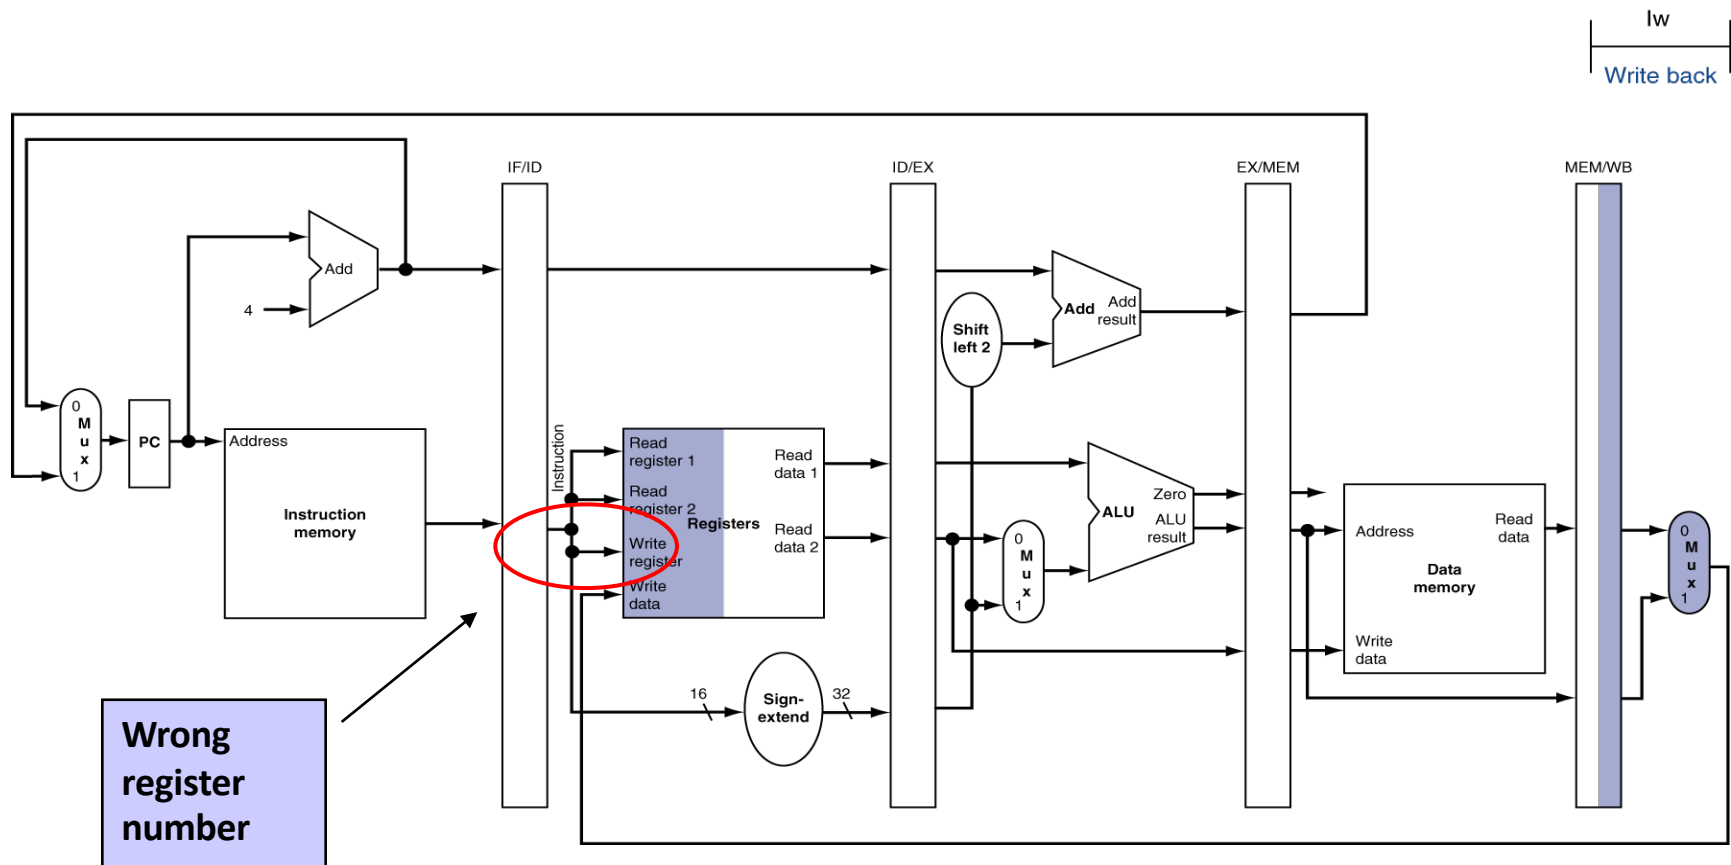
Classic 5-stage Pipeline for RISC: A First Shot

- Example: executing load (lw) instruction – **MEM stage**



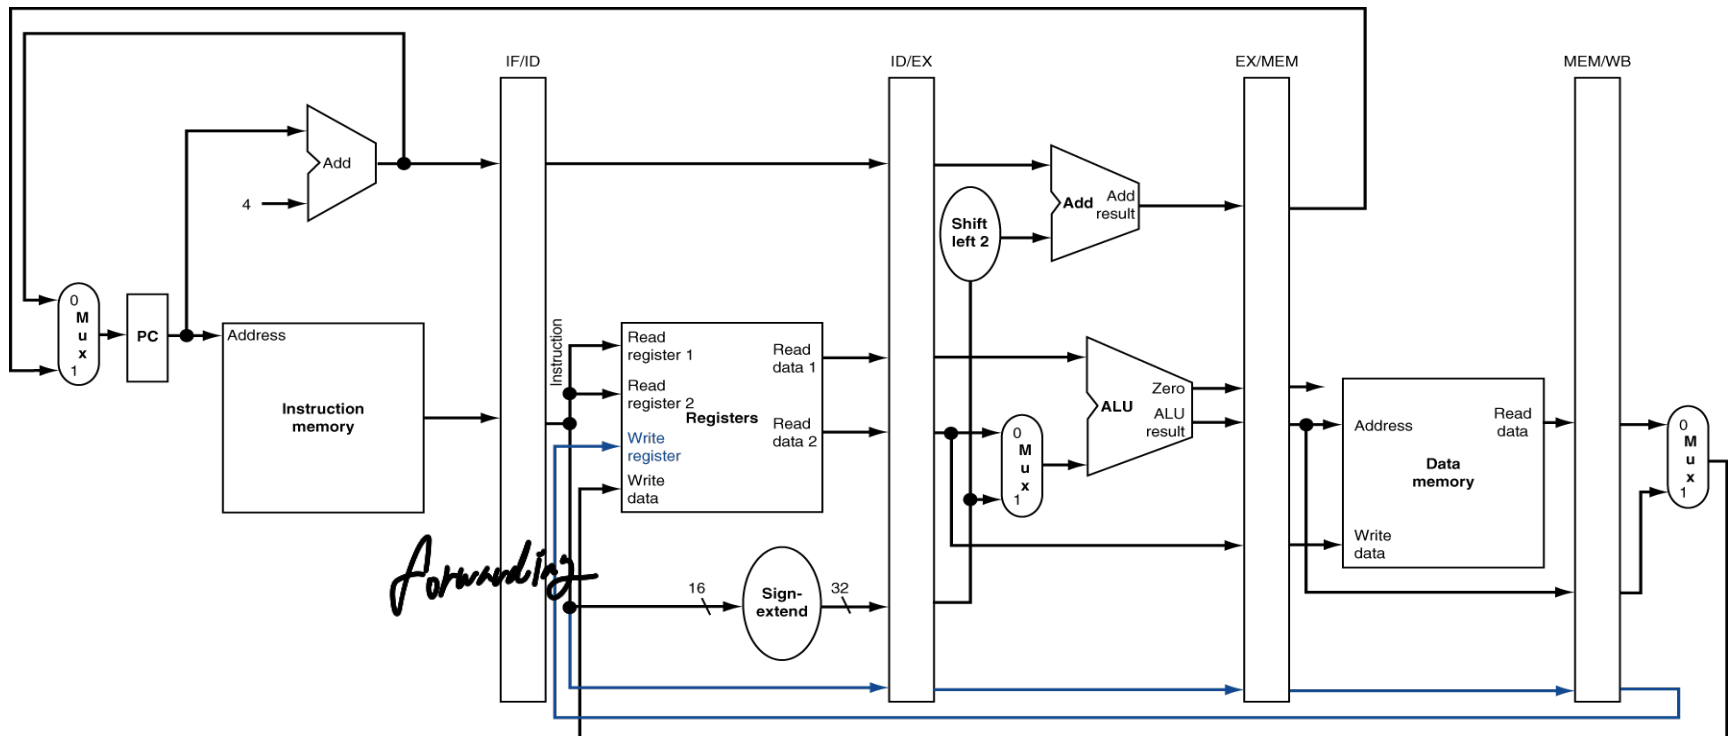
Classic 5-stage Pipeline for RISC: A First Shot

- Example: executing load (lw) instruction – **WB stage**



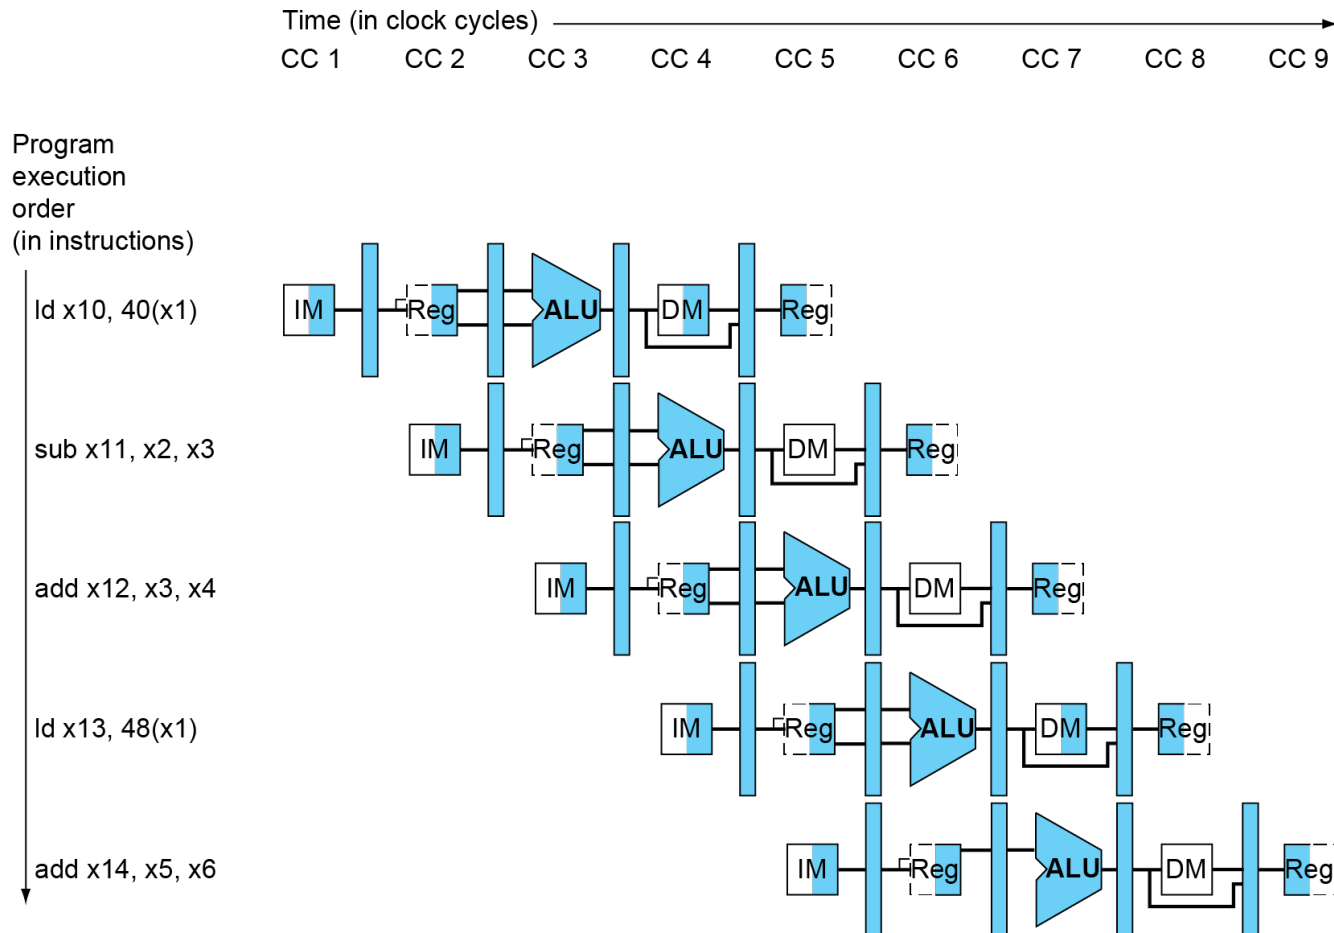
Classic 5-stage Pipeline for RISC: A First Shot

- Example: executing load (lw) instruction – **WB stage**
- **With corrected datapath**



Classic 5-stage Pipeline for RISC: A First Shot

■ Visualizing pipelining: showing resource usage

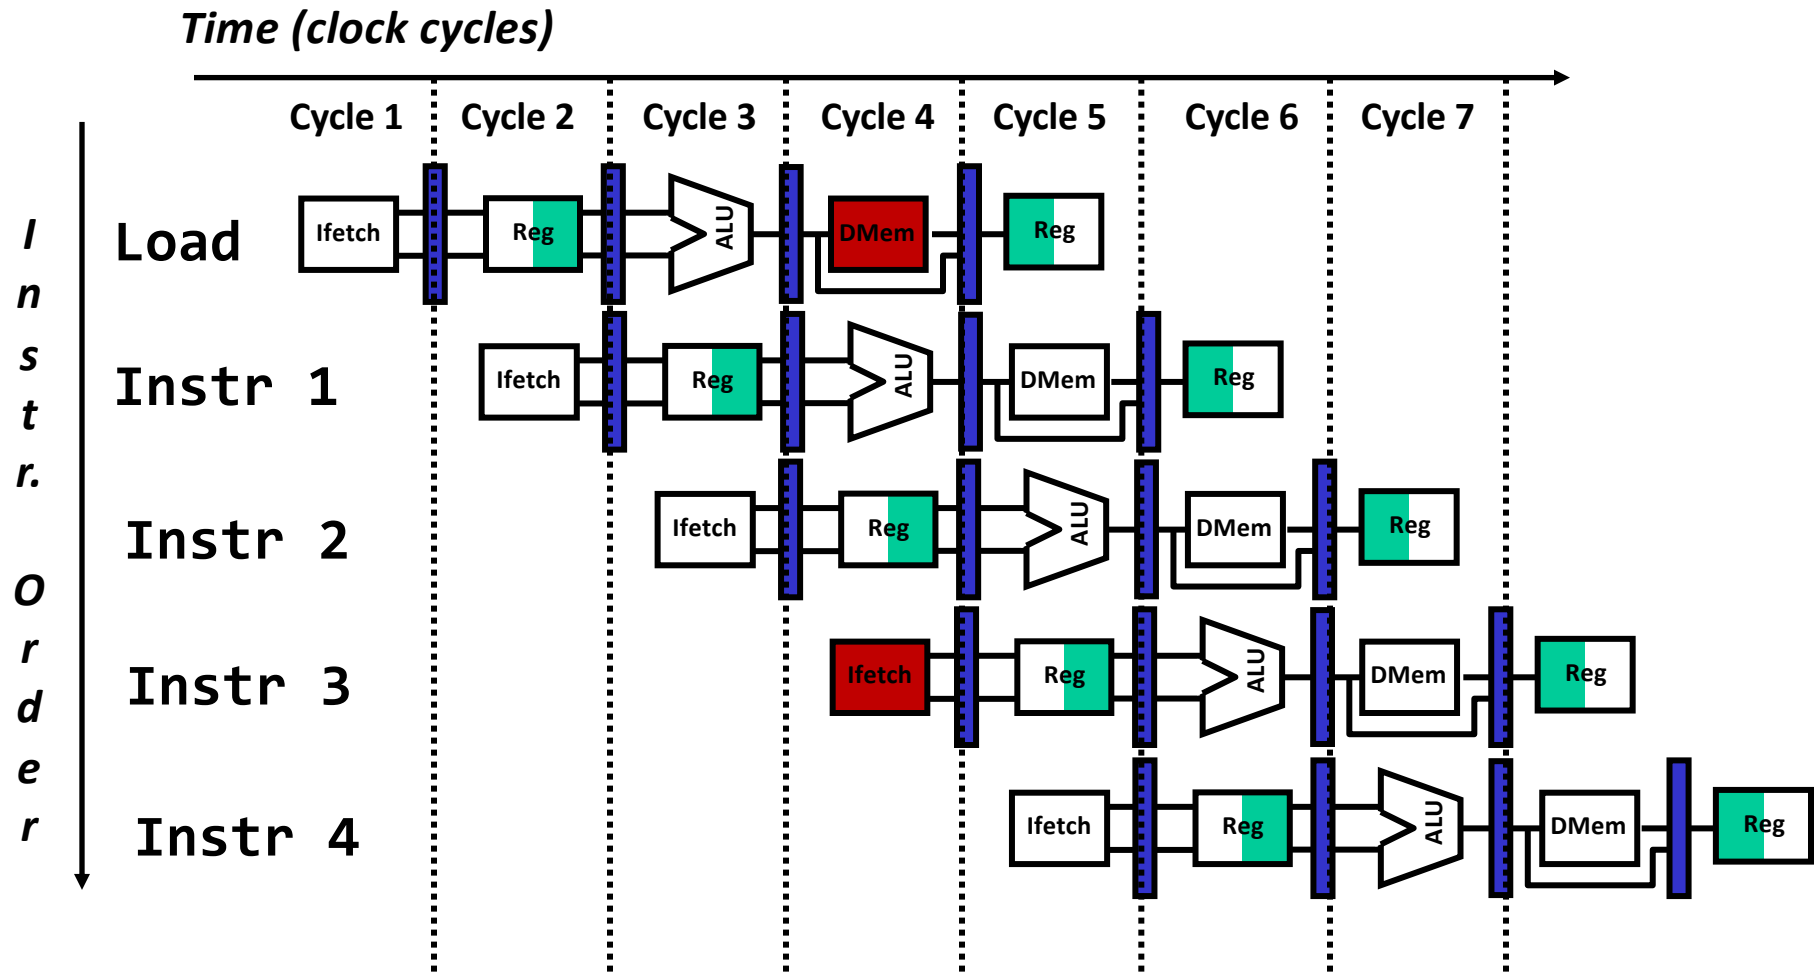


Pipeline Hazards

- Pipelining is not quite that easy!
- Limits to pipelining: **Hazards** prevent next instruction from executing during its designated clock cycle
 - Structural hazards: HW cannot support this combination of instructions
 - Data hazards: Instruction depends on result of prior instruction still in the pipeline
 - Control hazards: Caused by delay between the fetching of instructions and decisions about changes in control flow (branches and jumps).


Pipeline Hazards: Structural Hazards

■ Structural hazard with only one memory port



Pipeline Hazards: Structural Hazards


■ Structural hazard with only one memory port



Pipeline Hazards: Structural Hazards

■ Speedup equation for pipelining

- CPU time = Instruction_count x CPI x clock_cycle



$$CPI_{\text{pipelined}} = \text{Ideal CPI} + \text{Average Stall cycles per Inst}$$

$$\text{Speedup} = \frac{\text{Pipeline depth}}{\text{Ideal CPI} + \text{Pipeline stall CPI}} \times \frac{\text{Cycle Time}_{\text{unpipelined}}}{\text{Cycle Time}_{\text{pipelined}}}$$

For simple RISC pipeline, CPI = 1:

$$\text{Speedup} = \frac{\text{Pipeline depth}}{1 + \text{Pipeline stall CPI}} \times \frac{\text{Cycle Time}_{\text{unpipelined}}}{\text{Cycle Time}_{\text{pipelined}}}$$

pipeline stage ↑: more register,

Pipeline Hazards: Structural Hazards

■ Example: dual-port vs. single-port

- Machine A: Dual ported memory (“Harvard Architecture”)
- Machine B: Single ported memory (“Princeton Architecture”), but its pipelined implementation has a 1.05x faster clock rate
- Ideal CPI = 1 for both
- Suppose that Loads are 40% of instructions executed

$$\begin{aligned}\text{SpeedUp}_A &= \text{Pipeline Depth} / (1 + 0) \times (\text{clock}_{\text{unpipe}} / \text{clock}_{\text{pipe}}) \\ &= \text{Pipeline Depth}\end{aligned}$$

$$\begin{aligned}\text{SpeedUp}_B &= \text{Pipeline Depth} / (1 + 0.4 \times 1) \times (\text{clock}_{\text{unpipe}} / (\text{clock}_{\text{unpipe}} / 1.05)) \\ &= (\text{Pipeline Depth} / 1.4) \times 1.05 \\ &= 0.75 \times \text{Pipeline Depth}\end{aligned}$$

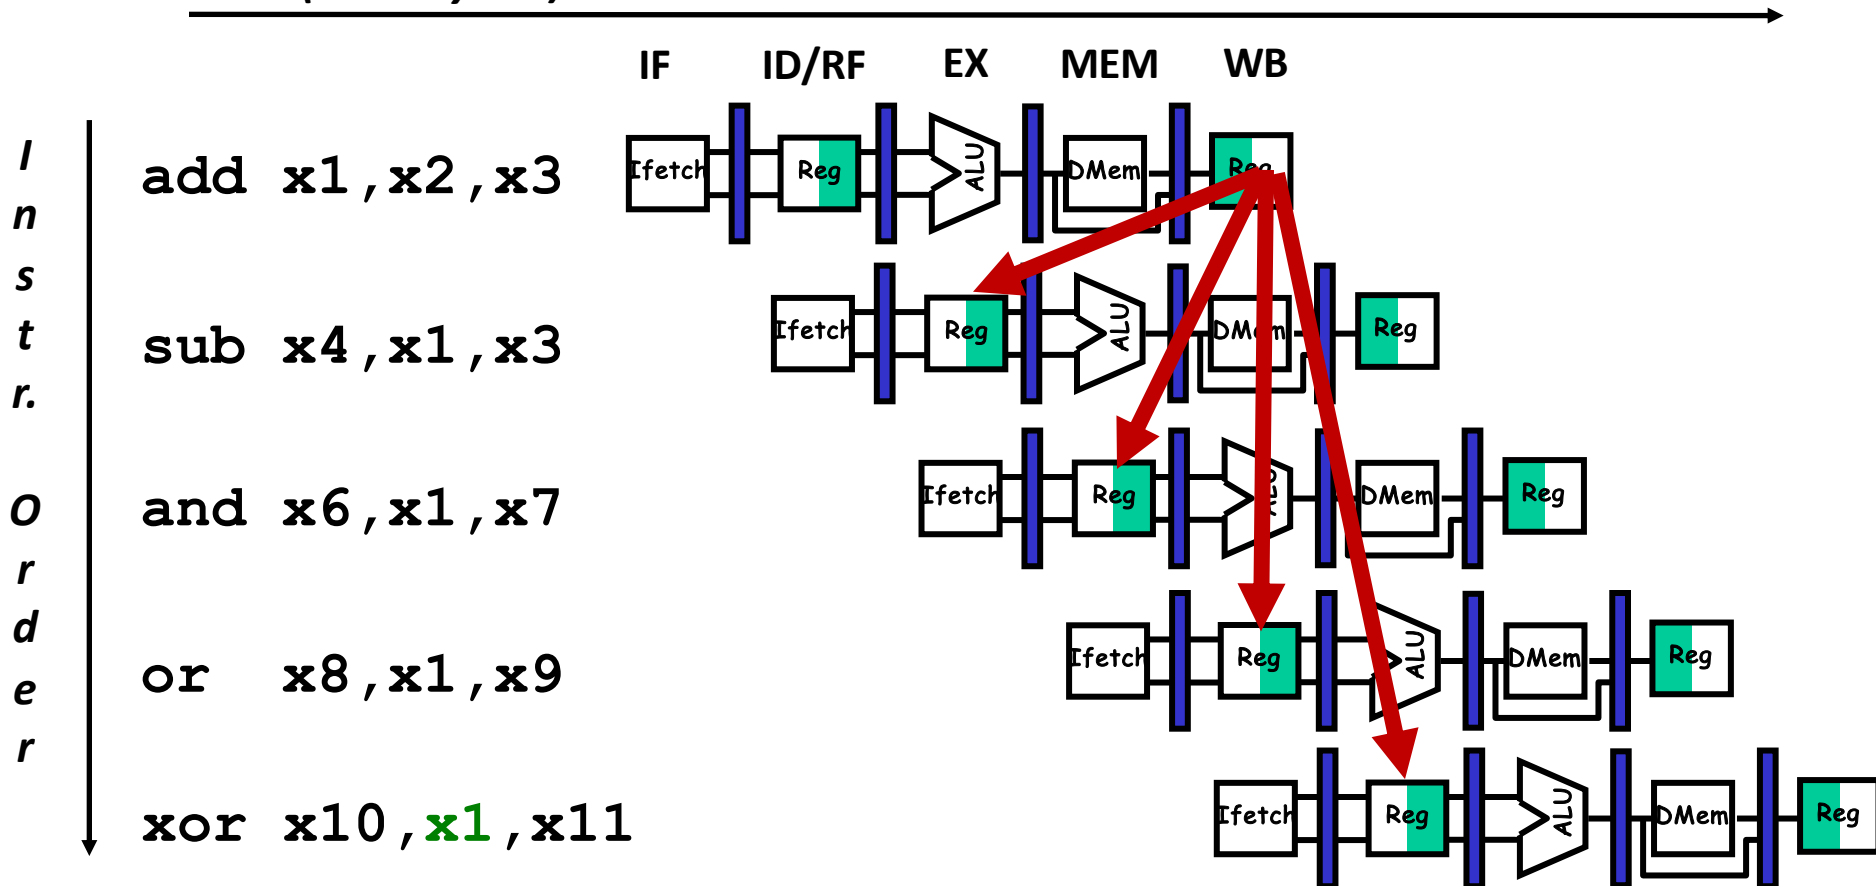
$$\text{SpeedUp}_A / \text{SpeedUp}_B = \text{Pipeline Depth} / (0.75 \times \text{Pipeline Depth}) = 1.33$$

- Machine A is 1.33x faster

Pipeline Hazards: Data Hazards

■ Data hazard on R1

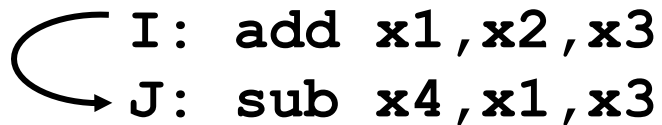
Time (clock cycles)



Pipeline Hazards: Data Hazards

■ Generic data hazards 1: **Read After Write (RAW)**

- Instr_j tries to read operand before Instr_i writes it



 I: add x1, x2, x3

 J: sub x4, x1, x3

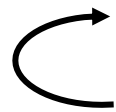
- Caused by a “**Dependence**” (in compiler nomenclature). This hazard results from an actual need for communication.

Hazard → dependency
 potential crisis (but not clear)
 ① far enough

Pipeline Hazards: Data Hazards

■ Generic data hazards 2: **Write After Read (WAR)**

- Instr_j writes operand before Instr_i reads it



```

I:  sub  x4, x1, x3
J:  add  x1, x2, x3
K:  mul  x6, x1, x7
  
```

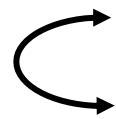
- Called an “**anti-dependence**” by compiler writers. This results from reuse of the name “**x1**”.
- Can’t happen in our RISC 5-stage pipeline because:
 - All instructions take 5 stages, and
 - Reads are always in stage 2, and
 - Writes are always in stage 5

Sol: register renaming?

Pipeline Hazards: Data Hazards

■ Generic data hazards 3: **Write After Write (WAW)**

- Instr_j writes operand before Instr_i writes it.



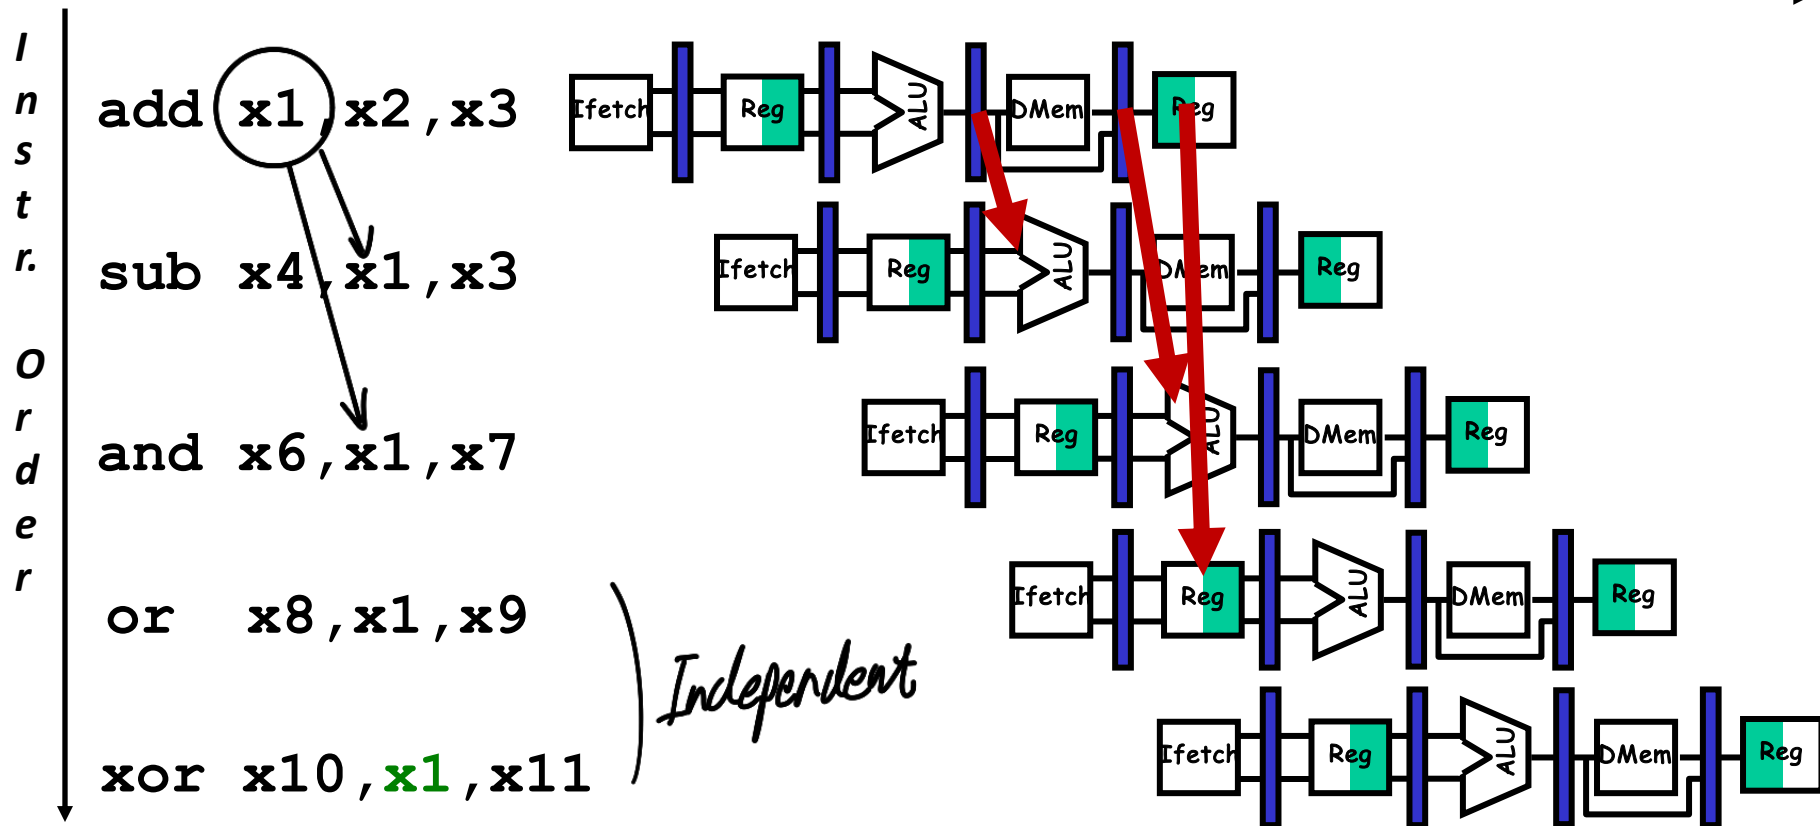
```
I: sub x1, x4, x3
J: add x1, x2, x3
K: mul x6, x1, x7
```

- Called an “**output dependence**” by compiler writers. This also results from the reuse of name “**x1**”.
- Can’t happen in our RISC 5 stage pipeline because:
 - All instructions take 5 stages, and
 - Writes are always in stage 5
- Will see WAR and WAW in more complicated pipes

Pipeline Hazards: Data Hazards

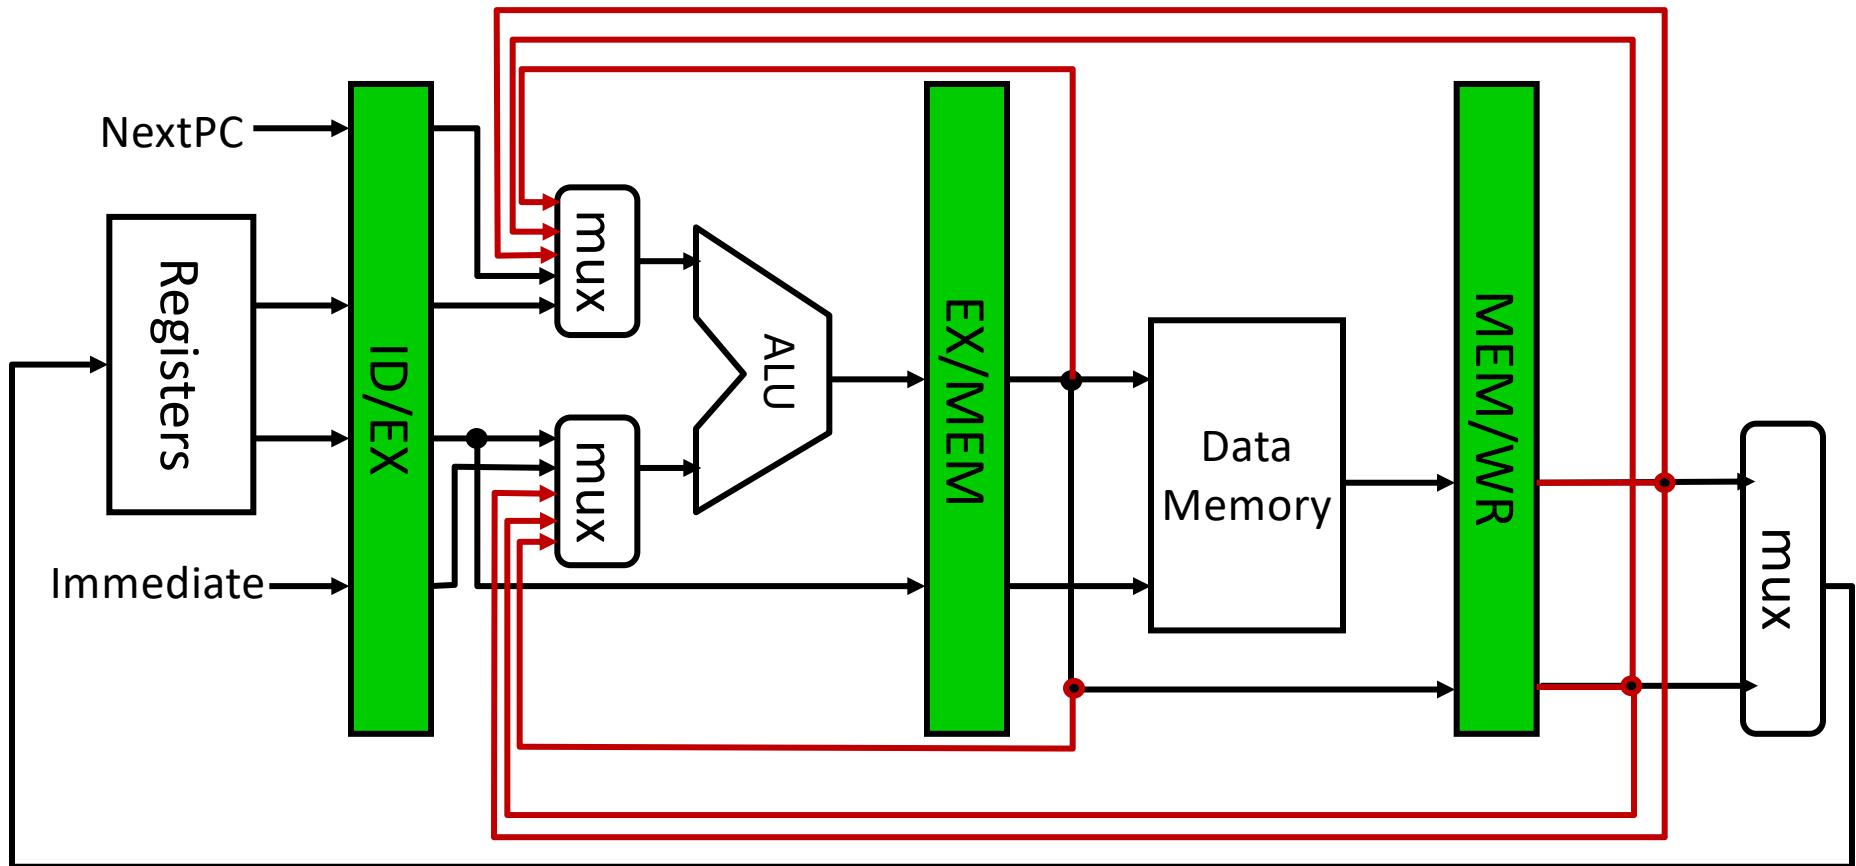
■ Forwarding to avoid data hazard

Time (clock cycles)



Pipeline Hazards: Data Hazards

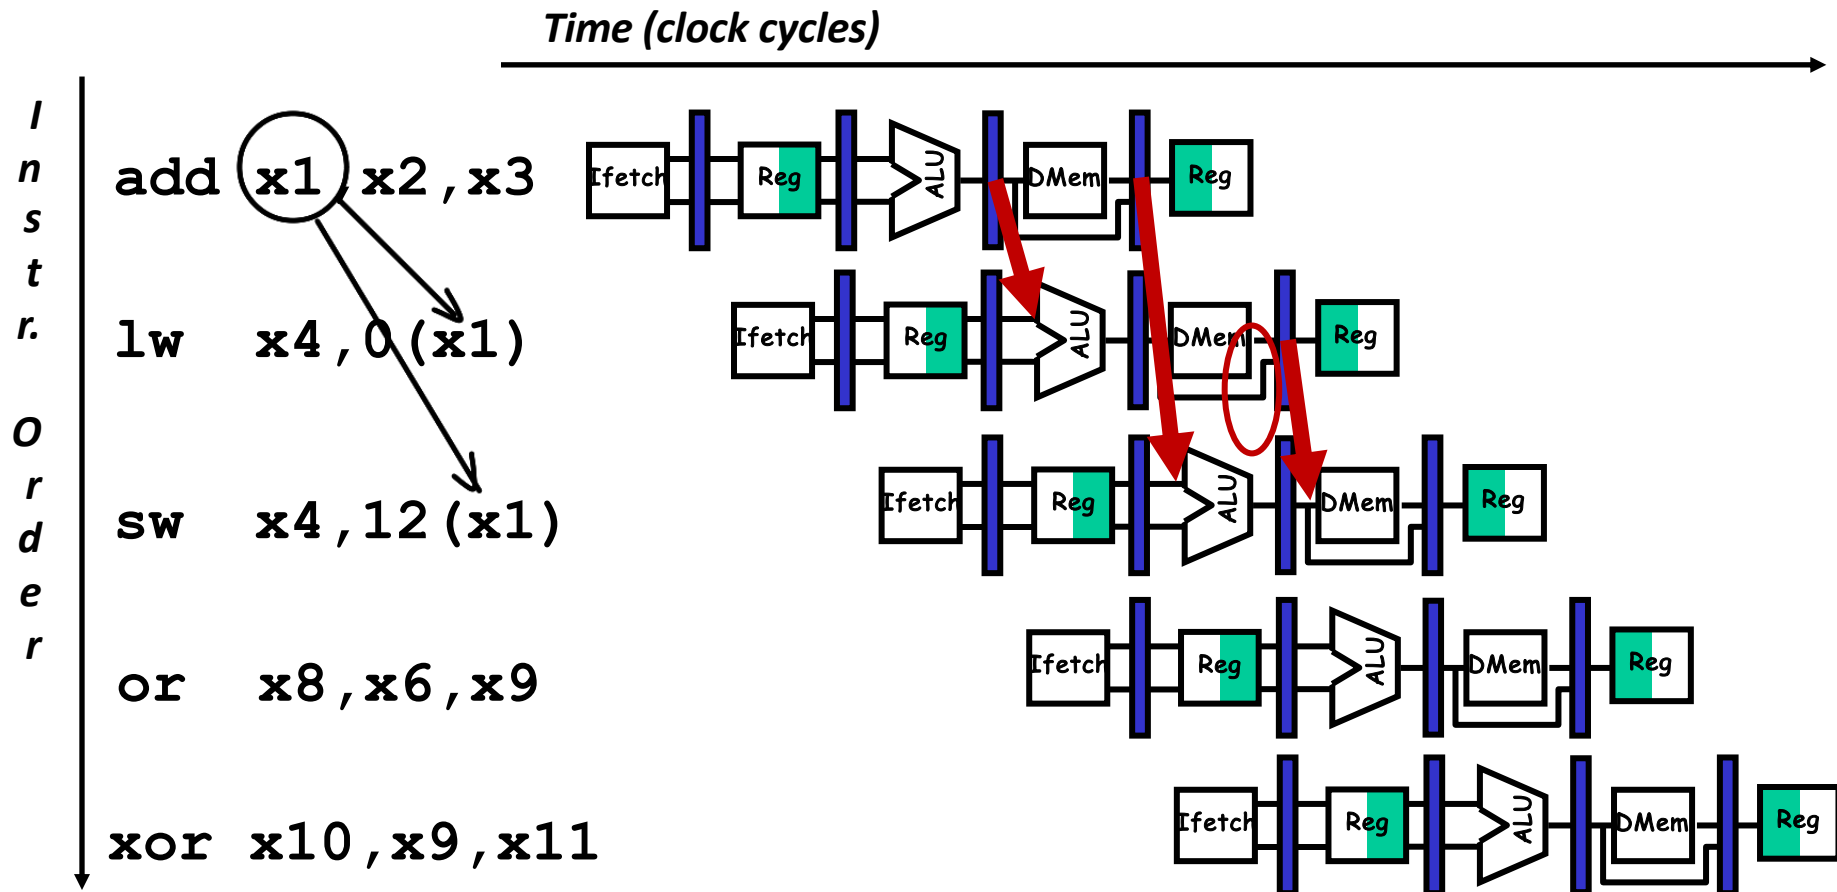
■ HW change for forwarding



What circuit detects and resolves this hazard?

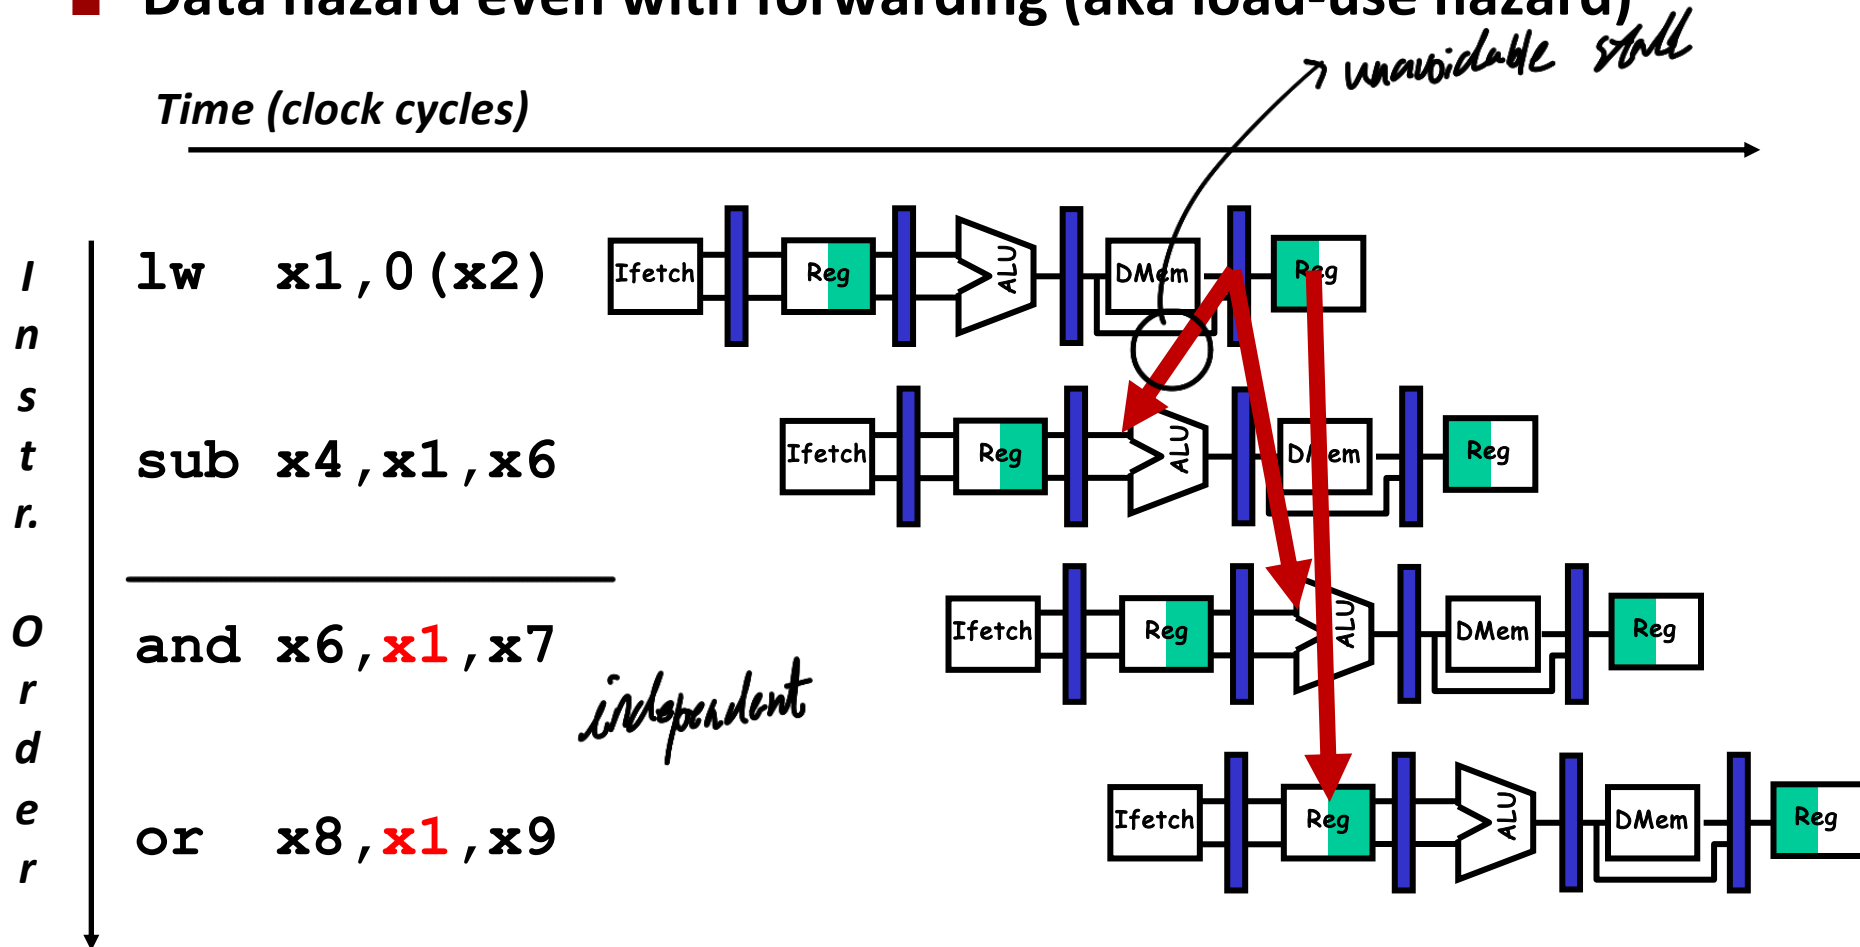
Pipeline Hazards: Data Hazards

■ Forwarding to avoid LW-SW data hazard



Pipeline Hazards: Data Hazards

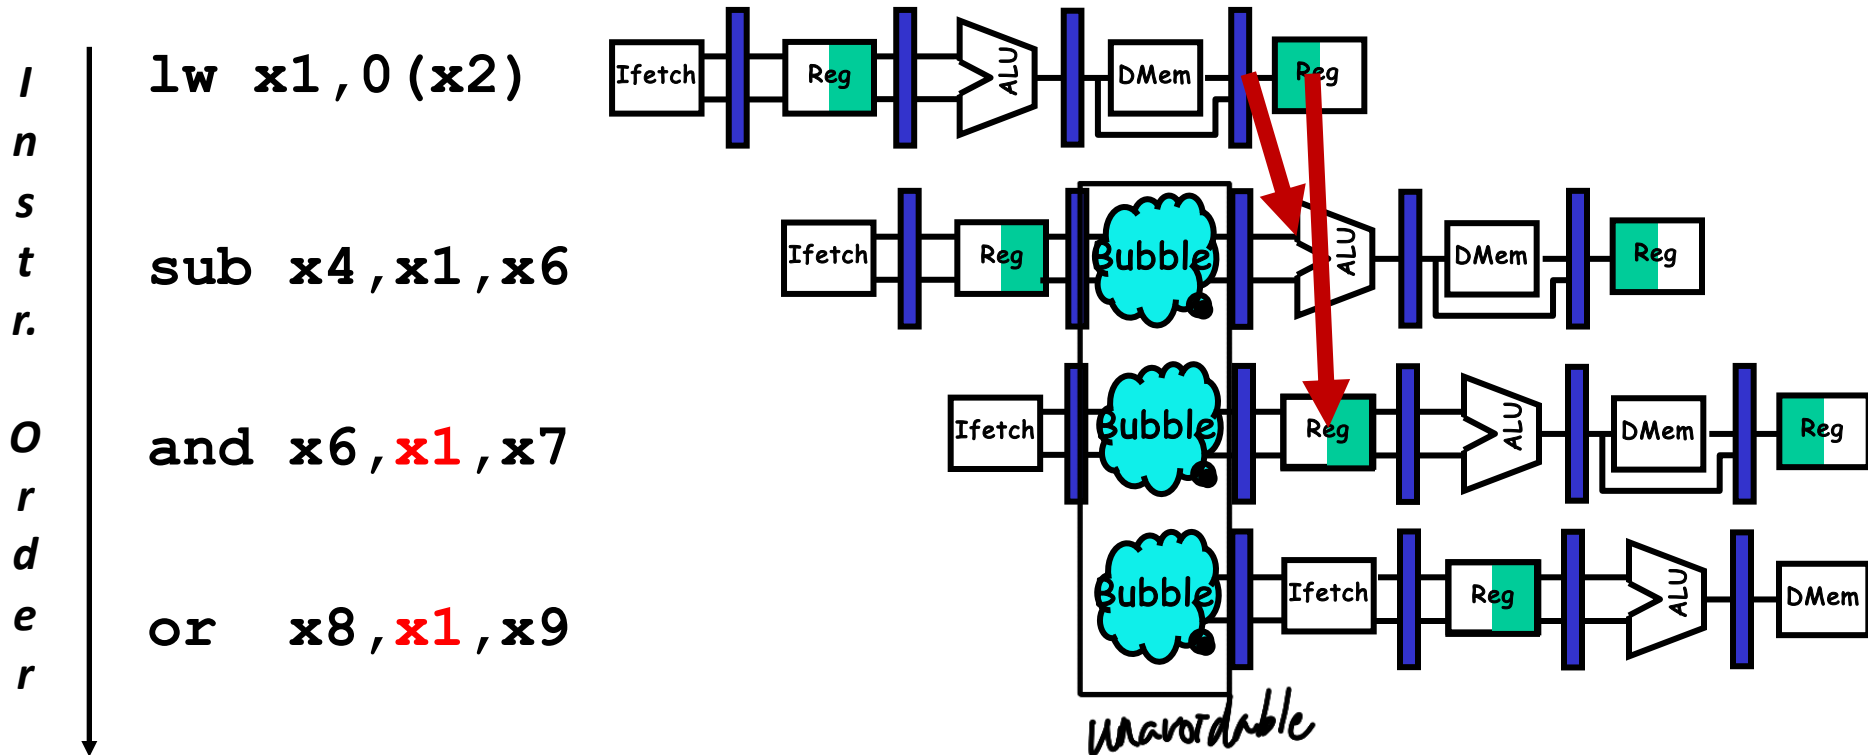
- Data hazard even with forwarding (aka load-use hazard)



Pipeline Hazards: Data Hazards

- Data hazard even with forwarding (aka load-use hazard)

Time (clock cycles)

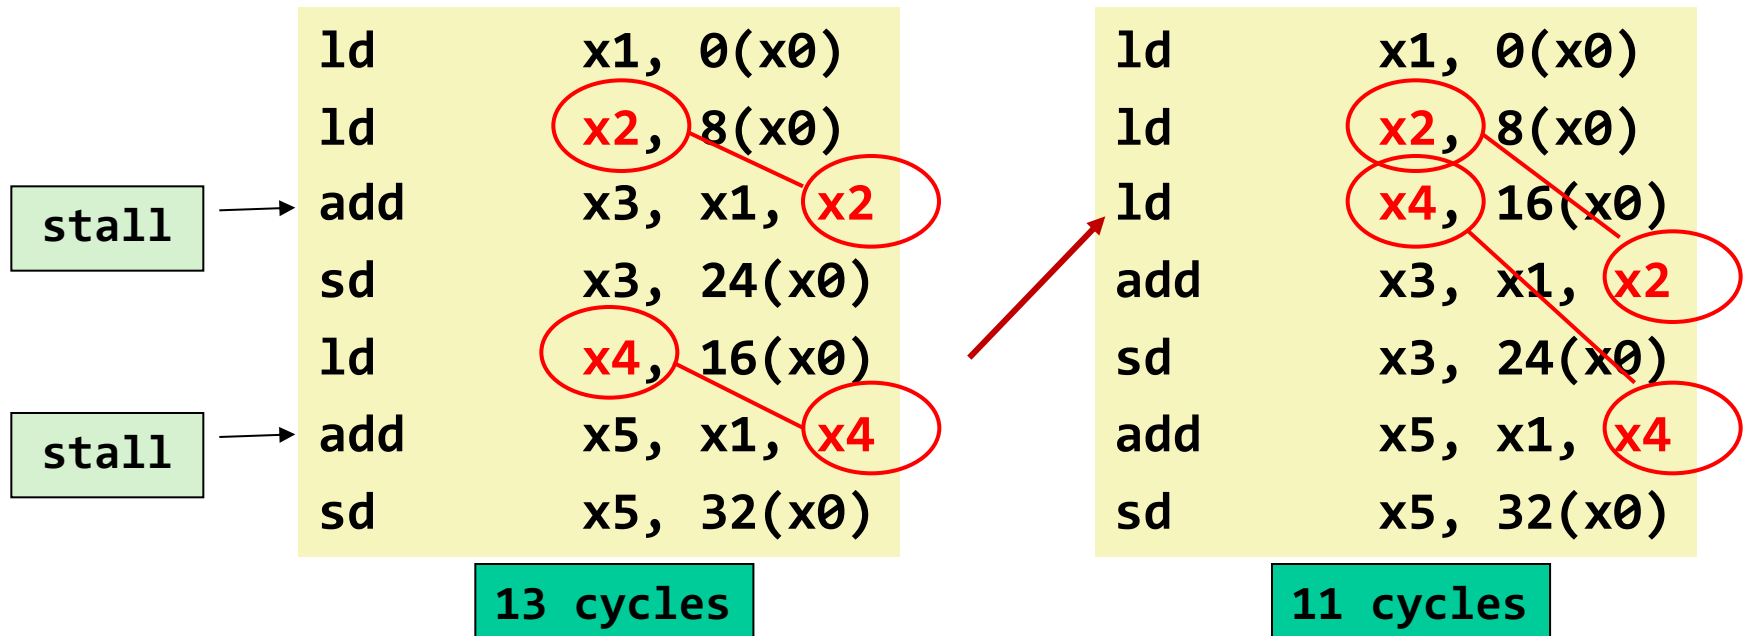


Pipeline Hazards: Data Hazards

- Software scheduling to avoid load hazards
 - Reorder code to avoid use of load result in the next instruction

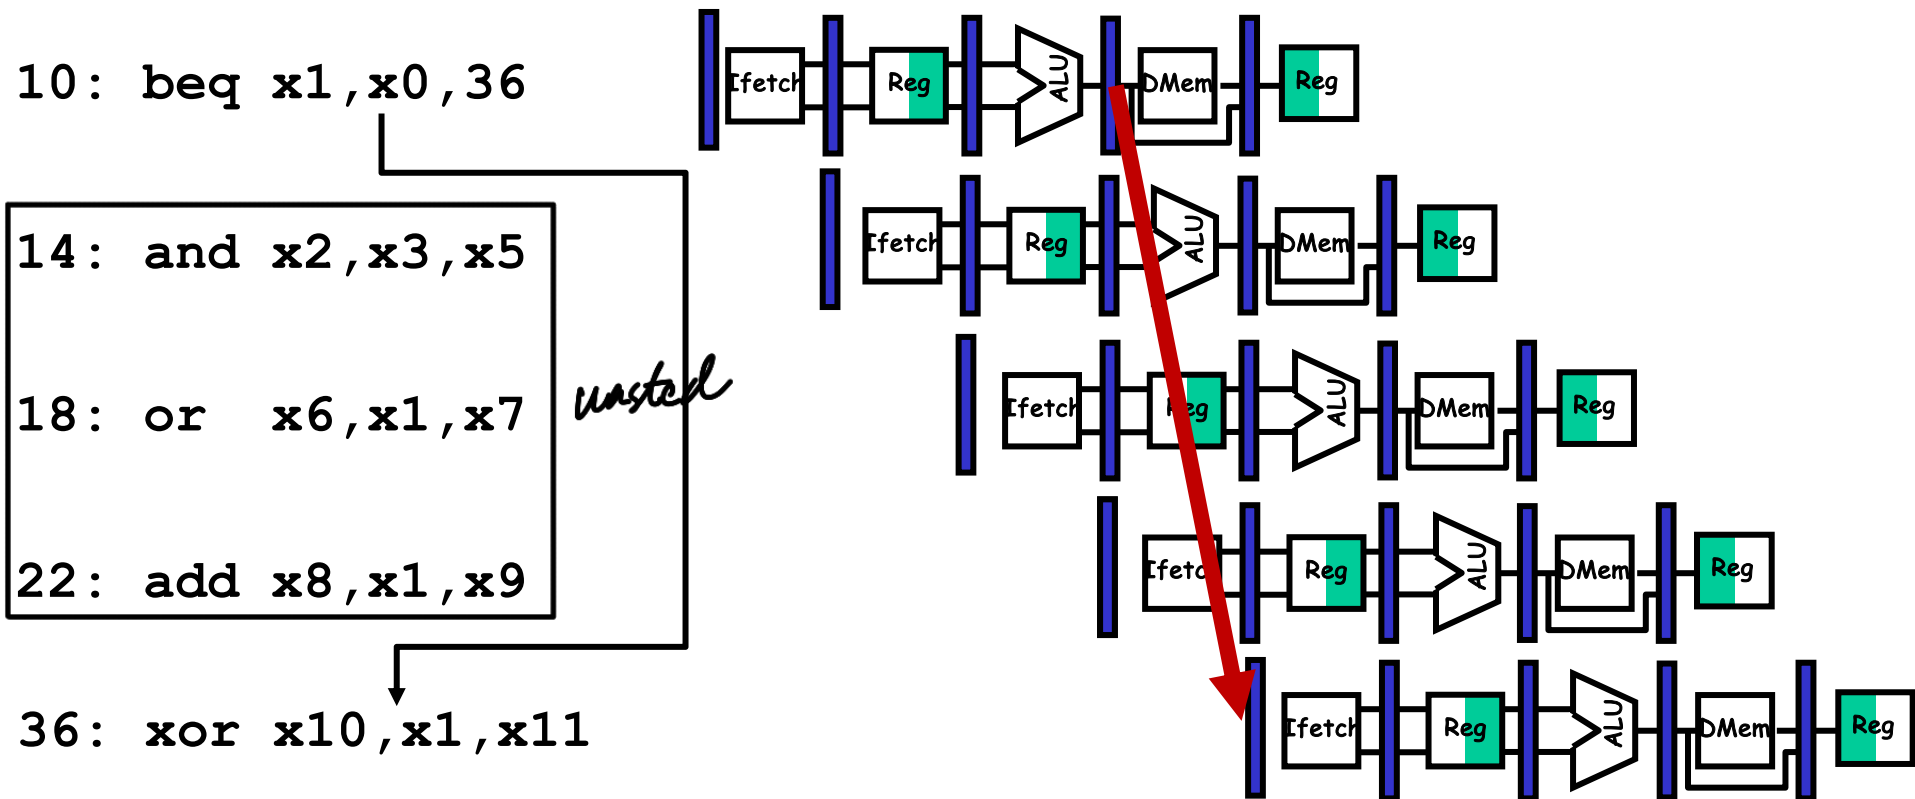
a = b + e;

d = b + f;



Pipeline Hazards: Control Hazards

■ Control hazard on branches: three-stage stall



What do you do with the 3 instructions in between?

How do you do it?

Where is the "commit"?

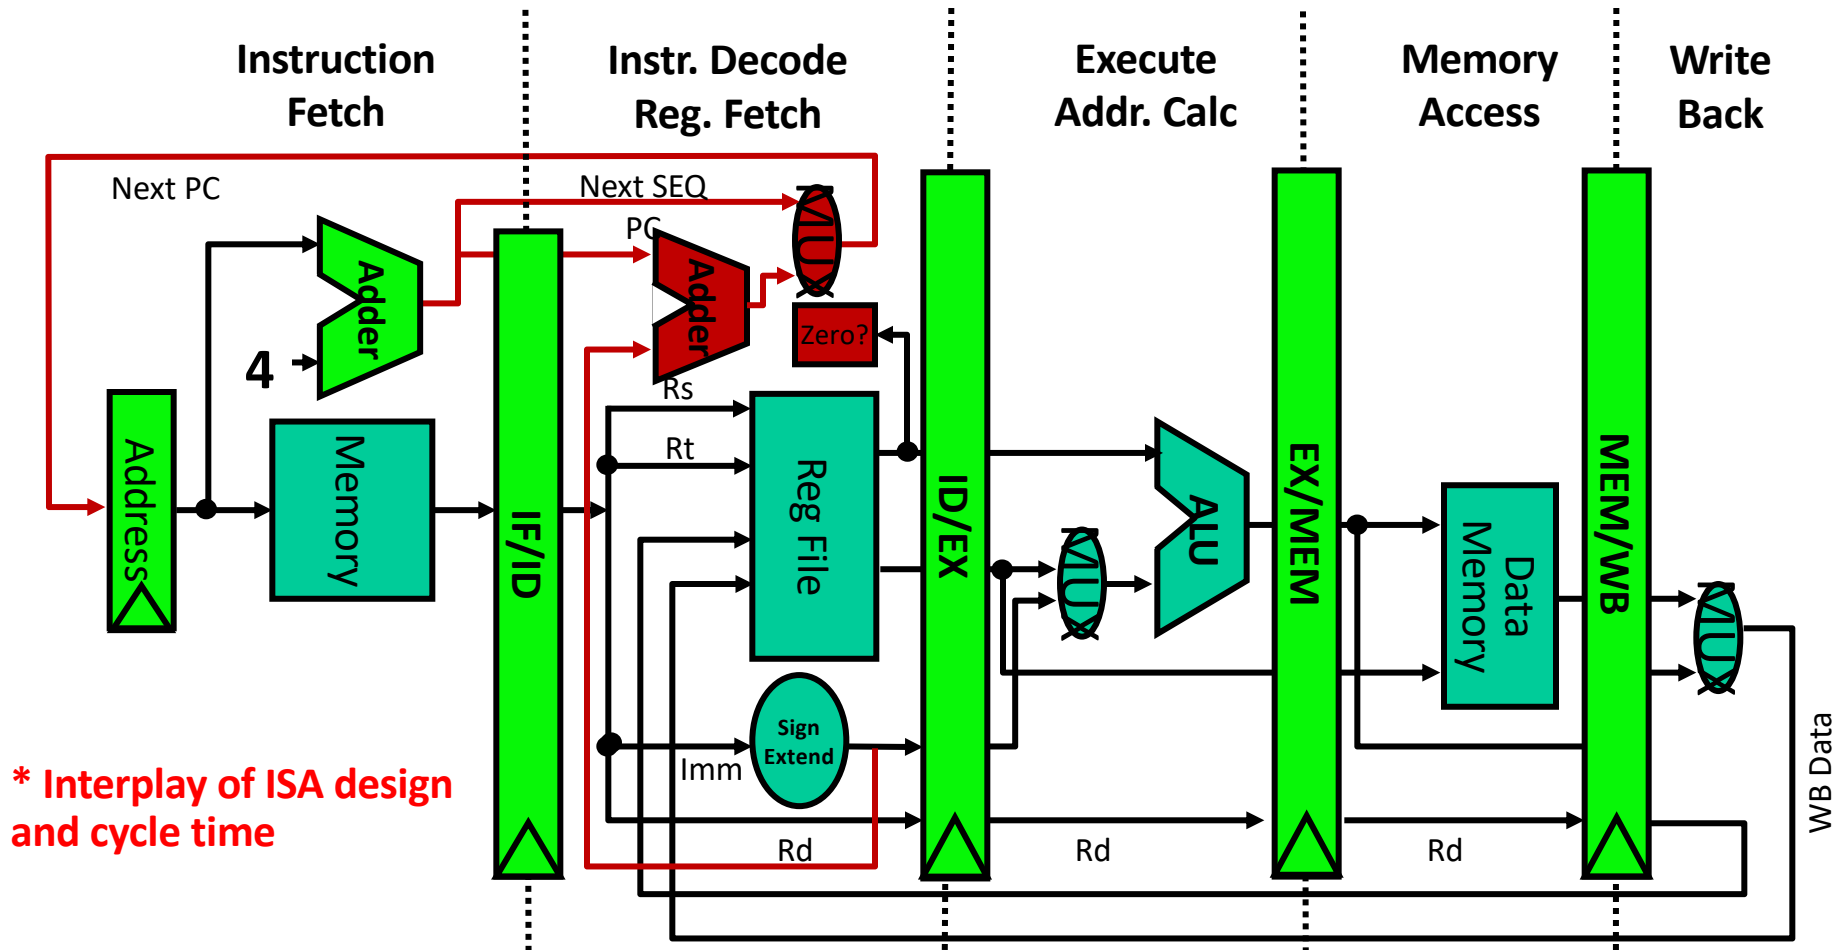
Pipeline Hazards: Control Hazards

■ Performance impact of branch stalls

- If CPI = 1, 30% branch, stall 3 cycles
 - New CPI = **1.9!**
- Two part solution:
 - Determine branch taken or not sooner, AND : *prediction*
 - Compute taken branch address earlier ; *add compare logic at ID stage*
- Branch tests if register = 0 or $\neq 0$
- Solution:
 - Move Zero test to ID/RF stage
 - Adder to calculate new PC in ID/RF stage
 - 1 clock cycle penalty for branch (versus 3)

Pipeline Hazards: Control Hazards

■ Pipelined RISC datapath (new)



Pipeline Hazards: Control Hazards

■ Four branch hazard alternatives

#1: Stall until branch direction is clear

#2: Predict Branch Not Taken

- Execute successor instructions in sequence
- “Squash” instructions in pipeline if branch actually taken
- Advantage of late pipeline state update
- 47% branches not taken on average (MIPS)
- PC+4 already calculated, so use it to get next instruction

#3: Predict Branch Taken

- 53% branches taken on average (MIPS)
- But haven't calculated branch target address in MIPS
 - It still incurs 1 cycle branch penalty – no performance benefit
 - Other machines: branch target known before outcome

Pipeline Hazards: Control Hazards

■ Four branch hazard alternatives (cont)

#4: Delayed Branch

- Define branch to take place **AFTER** a following instruction

branch instruction


sequential successor₁

sequential successor₂

.....

sequential successor_n

branch target if taken

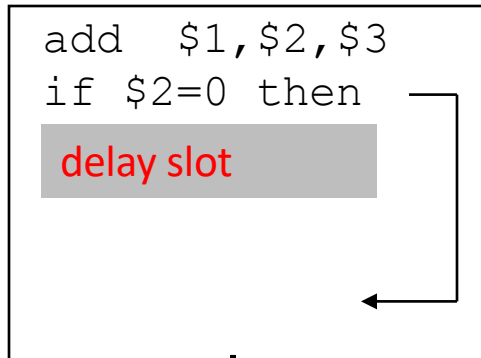


Branch delay of length n


- 1 slot delay allows proper decision and branch target address in 5 stage pipeline
- MIPS uses this

Pipeline Hazards: Control Hazards

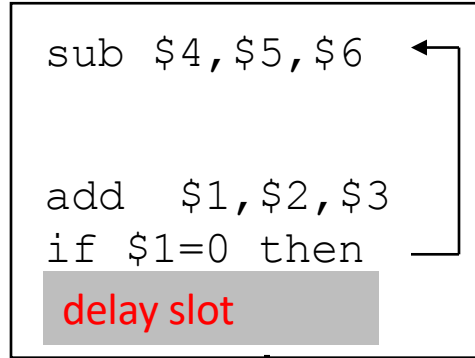
A. From before branch




becomes



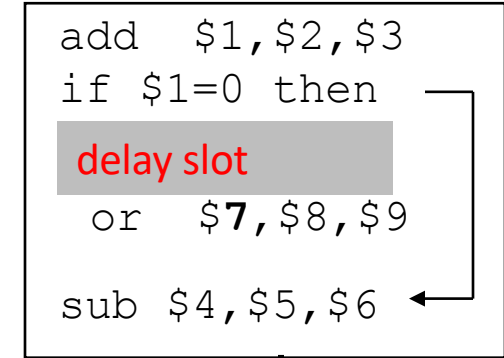
B. From branch target



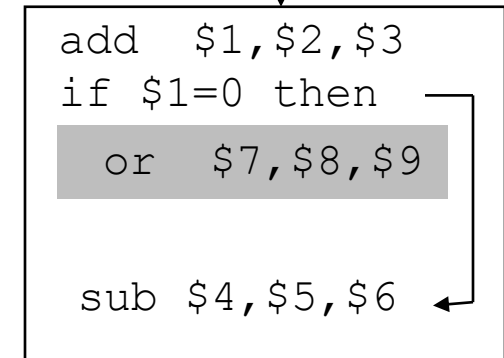
becomes



C. From fall through



becomes



■ Instruction scheduling branch delay slots

- A is the best choice, fills delay slot & reduces instruction count (IC)
- In B, the `sub` instruction may need to be copied, increasing IC
- In B and C, must be okay to execute ~~sub~~ `sub` or `or` when branch fails

Pipeline Hazards: Control Hazards

■ Delayed branch

- Compiler effectiveness for single branch delay slot:
 - Fills about 60% of branch delay slots
 - About 80% of instructions executed in branch delay slots useful in computation
 - About 50% ($60\% \times 80\%$) of slots usefully filled
- Delayed Branch downside: As processor go to deeper pipelines and multiple issue, the branch delay grows and need more than one delay slot
 - Delayed branching has lost popularity compared to more expensive but more flexible dynamic approaches
 - Growth in available transistors has made dynamic approaches relatively cheaper

Branch prediction

Exception/Interrupt Handling

■ Exceptions vs. interrupts

internal

- **Exception:** An unusual event happens to an instruction during its execution
 - Examples: divide by zero, undefined opcode

- **Interrupt:** Hardware signal to switch the processor to a new instruction stream

external

What?
✓

- Example: a sound card interrupts when it needs more audio output samples (an audio “click” happens if it is left waiting)

■ “Precise” exceptions

- Problem: It must appear that the exception or interrupt must appear between 2 instructions (I_i and I_{i+1})
 - The effect of all instructions up to and including I_i is complete
 - No effect of any instruction after I_i can take place
- The interrupt (exception) handler either aborts program or restarts at instruction I_{i+1}

Exception/Interrupt Handling

■ Exception handling in RISC-V

- Save PC of offending (or interrupted) instruction
 - In RISC-V: Supervisor Exception Program Counter (SEPC)
- Save indication of the problem
 - In RISC-V: Supervisor Exception Cause Register (SCAUSE)
 - 64 bits, but most bits unused
 - Exception code field: 2 for undefined opcode, 12 for hardware malfunction, ...
- Jump to handler (*inside OS*)
 - Assume at 0000 0000 1C09 0000_{hex}

return PC
↓
reason

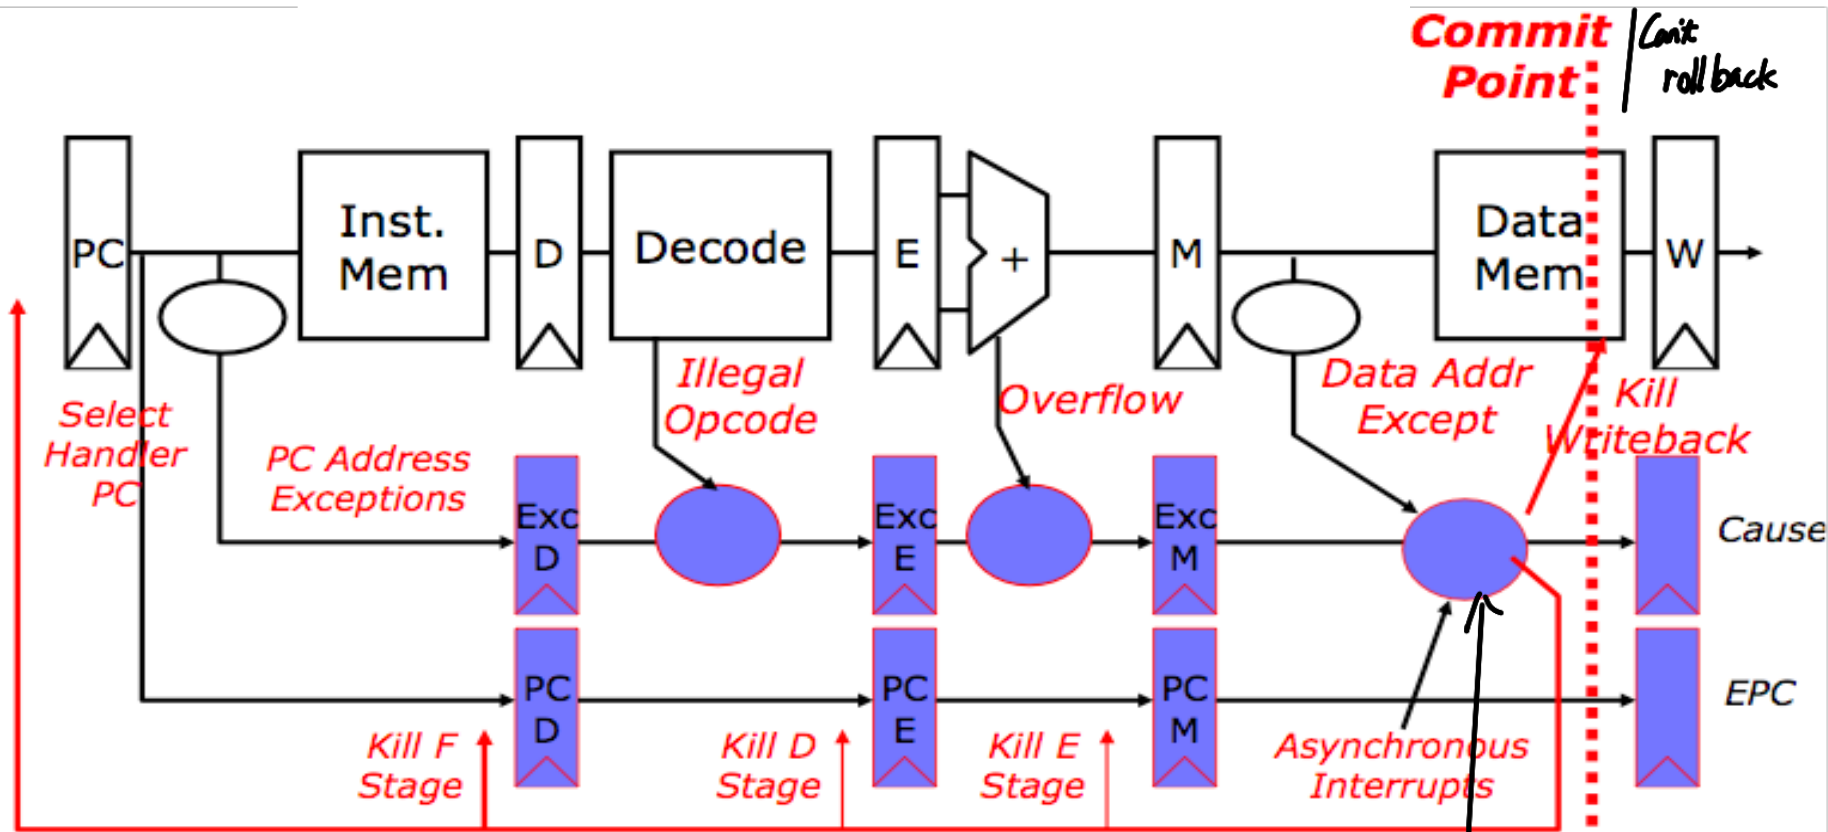
Exception/Interrupt Handling

■ Handler Actions

- Read cause, and transfer to relevant handler
- Determine action required
- If restartable
 - Take corrective action
 - use SEPC to return to program
- Otherwise
 - Terminate program
 - Report error using SEPC, SCAUSE, ...

Exception/Interrupt Handling

■ Precise exceptions in static pipelines



Key observation: architected state only change in memory and register write stages.

Priority

Exceptions in a Pipeline

- **Another form of control hazard**
- **Consider malfunction on add in EX stage**
add x1, x2, x1
 - Prevent x1 from being clobbered
 - Complete previous instructions
 - Flush add and subsequent instructions
 - Set SEPC and SCAUSE register values
 - Transfer control to handler
- **Similar to mispredicted branch**
 - Use much of the same hardware

Handling Multi-Cycle Operations

- **Pipelining becomes complex when we want high performance in the presence of:**
 - Long latency or partially pipelined floating-point units
 - Memory systems with variable access time
 - Multiple arithmetic and memory units

Handling Multi-Cycle Operations


■ Floating-Point Unit (FPU)

- Much more hardware than an integer unit
 - Single-cycle FPU is a bad idea – why?
- Common to have several FPU's
- Common to have different types of FPU's: Fadd, Fmul, Fdiv, ...
- An FPU may be pipelined, partially pipelined or not pipelined
- To operate several FPU's concurrently the FP register file needs to have more read and write ports

Handling Multi-Cycle Operations

■ Functional Unit Characteristics

*fully
pipelined*



*partially
pipelined*



■ Functional units have internal pipeline registers

- Operands are latched when an instruction enters a functional unit
- Following instructions are able to write register file during a long-latency operation

Handling Multi-Cycle Operations

■ Floating-Point ISA

- Interaction between FP datapath and integer datapath is determined by ISA
- RISC-V ISA
 - Separate register files for FP and Integer instructions
 - the only interaction is via a set of move/convert instructions (some ISA's don't even permit this)
 - Separate load/store for FPR's and GPR's but both use GPR's for address calculation
 - Branches on FP values
 - Utilizing GPR for holding comparison results
 - feq (equal), flt (less than), fle (less than or equal to)

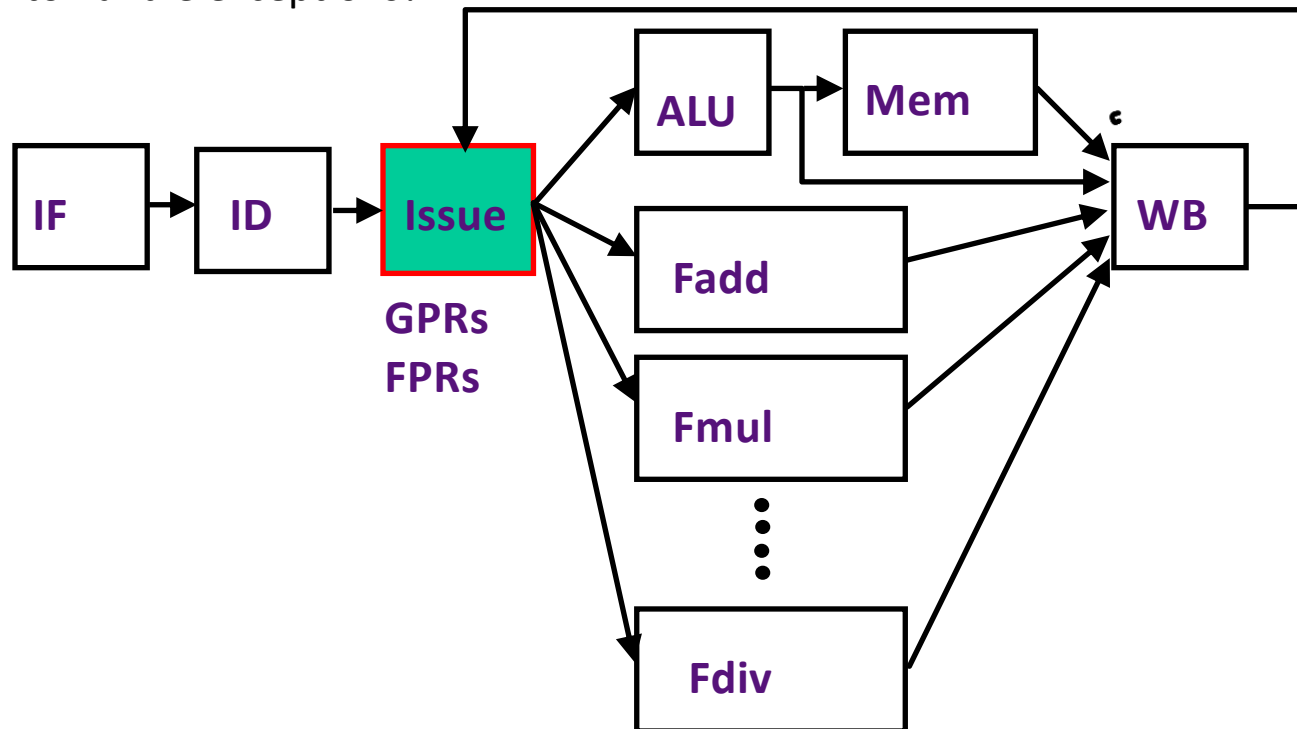
Handling Multi-Cycle Operations

- **Common approaches to improving memory performance in realistic memory systems:**
 - Caches - single cycle except in case of a miss \Rightarrow stall
 - Banked memory - multiple memory accesses \Rightarrow bank conflicts
 - split-phase memory operations (separate memory request from response), many in flight \Rightarrow out-of-order responses
- **Latency of access to the main memory is usually much greater than one cycle and often unpredictable**
 - Solving this problem is a central issue in computer architecture

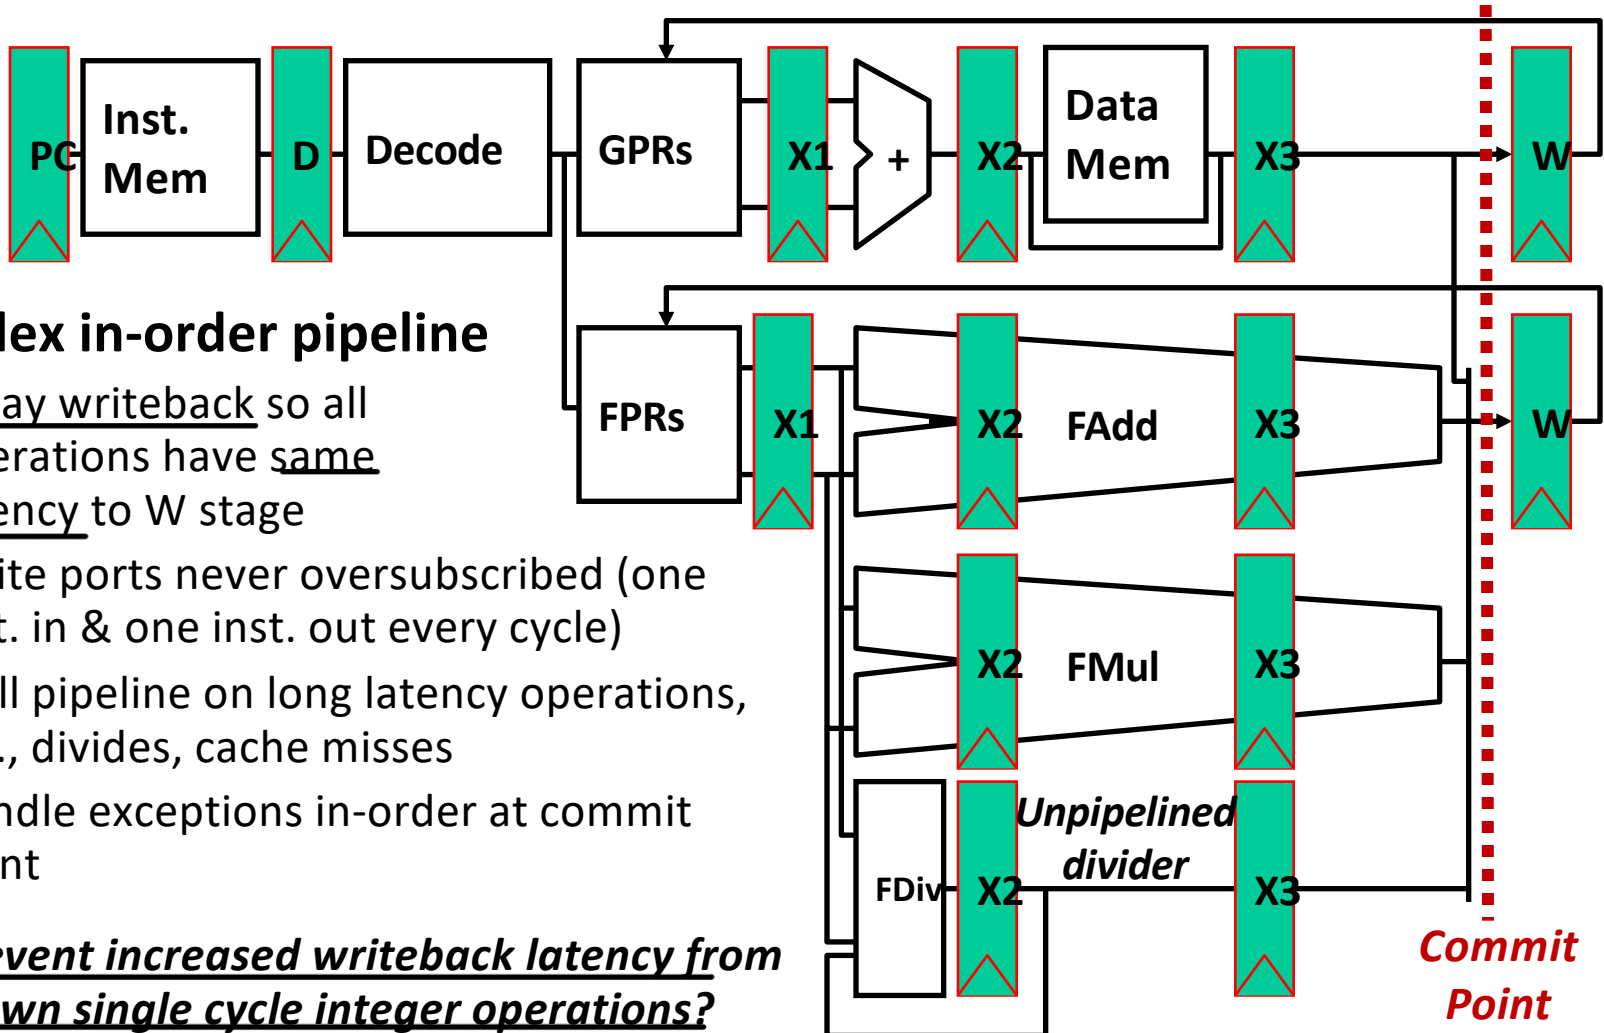
Handling Multi-Cycle Operations

■ Issues in Complex Pipeline Control

- Structural conflicts at the execution stage if some FPU or memory unit is not pipelined and takes more than one cycle
- Structural conflicts at the write-back stage due to variable latencies of different functional units
- Out-of-order write hazards due to variable latencies of different functional units
- How to handle exceptions?



Handling Multi-Cycle Operations



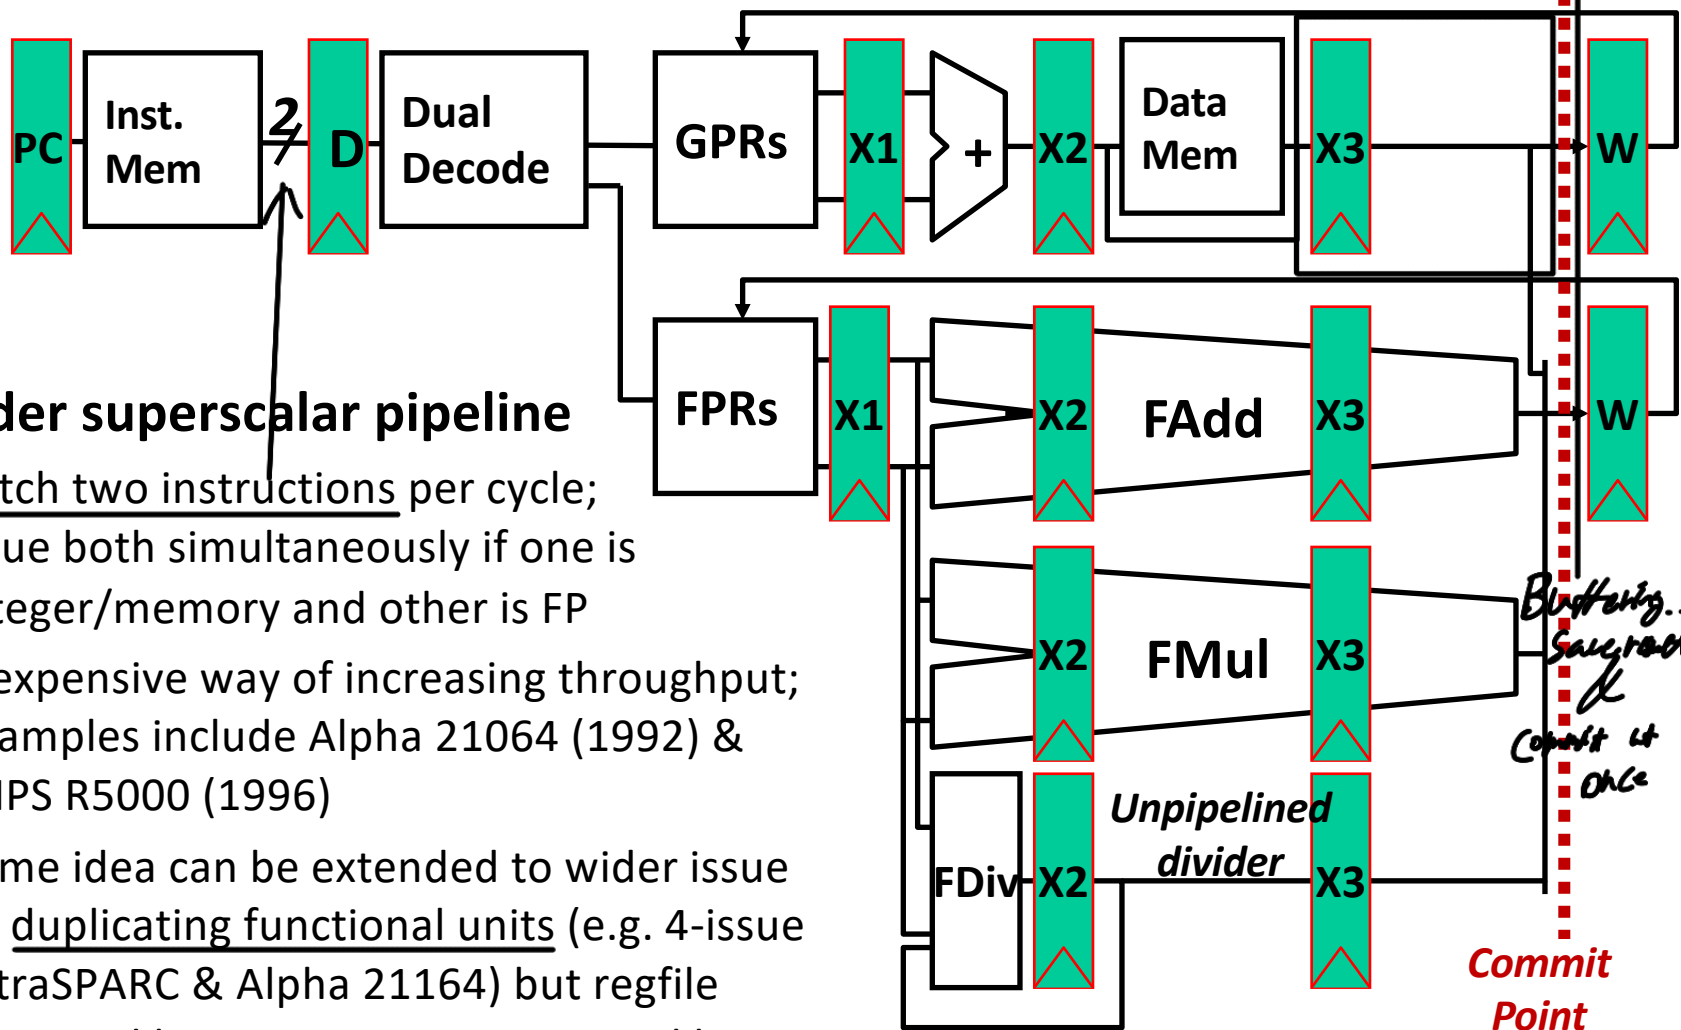
■ Complex in-order pipeline

- Delay writeback so all operations have same latency to W stage
- Write ports never oversubscribed (one inst. in & one inst. out every cycle)
- Stall pipeline on long latency operations, e.g., divides, cache misses
- Handle exceptions in-order at commit point

How to prevent increased writeback latency from slowing down single cycle integer operations?

Bypassing (=Forwarding)

Handling Multi-Cycle Operations



■ In-order superscalar pipeline

- Fetch two instructions per cycle; issue both simultaneously if one is integer/memory and other is FP
- Inexpensive way of increasing throughput; examples include Alpha 21064 (1992) & MIPS R5000 (1996)
- Same idea can be extended to wider issue by duplicating functional units (e.g. 4-issue UltraSPARC & Alpha 21164) but regfile ports and bypassing costs grow quickly

Summary

- Just overlap tasks; easy if tasks are independent
- Speed Up \leq Pipeline Depth; if ideal CPI is 1, then:

$$\text{Speedup} = \frac{\text{Pipeline depth}}{1 + \text{Pipeline stall CPI}} \times \frac{\text{Cycle Time}_{\text{unpipelined}}}{\text{Cycle Time}_{\text{pipelined}}}$$

- Hazards limit performance on computers:
 - Structural: need more HW resources
 - Data (RAW,WAR,WAW): need forwarding, compiler scheduling
 - Control: delayed branch, prediction
- Exceptions, interrupts add complexity
- Handling multi-cycle operations adds additional complexity